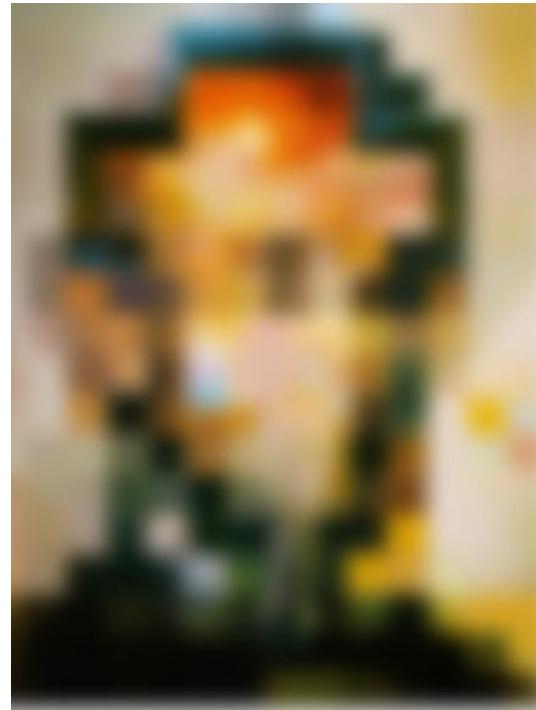




Computer Vision : Insight of object detection algorithms :

# Lesson1 Image Preprocessing



**By Daniel**  
Research Scientist  
[danielshaving@gmail.com](mailto:danielshaving@gmail.com)

## 课程大纲：

### 阶段1 图像预处理

#### 第一课：图像处理基础

知识点：图像处理，灰度值提取，Histogram处理，视觉和图像变换（Warp），图像卷积

#### 第二课：图像处理进阶：图像滤波，特征提取及匹配

知识点：Harris角点检测，Sift，Surf，Orb，图像特征点匹配，边缘检测算法等(含所有带工具箱代码 和 纯numpy代码)

#### 第三课：实践：利用KNN算法和OpenCV进行手写字符识别

### 阶段2：创建自己的图像识别神经网络

#### 第四课：深入理解神经网络的前向传递和反向传播及其物理意义

知识点：Loss function，交叉熵代价函数，梯度下降法求导

#### 第五课：训练你自己的网络，重点为调参和工作中用到的一些技巧

知识点：Loss function，交叉熵代价函数，梯度下降法

#### 第六课：卷积神经网络（RNN）在图像分类识别中的应用（附python编程和算法解析）

知识点：数据输入层，卷积计算层，激励层（Sigmoid, Tanh, ReLu, ELU），池化层，全连接层，Batch Normalization，学习率

#### 第七课：实践，不使用任何工具包，训练一个属于你自己的神经网络进行手写字符识别

### 系列3. 深度卷积神经网络进阶

#### 第八课：不同的神经网络类别和应用

知识点：调参基本技巧，向量点积

#### 第九课：深度卷积神经网络原理及实践

知识点：神经网络的迁移学习技巧

#### 第十课：搭建图片搜索系统，深入理解Triplet Loss 及其训练技巧

#### 第十一课：实践：使用Tensorflow/Keras搭建神经网络，进行图像分类

### 阶段4：目标检测和LSTM标注法

#### 第十二课：目标检测算法

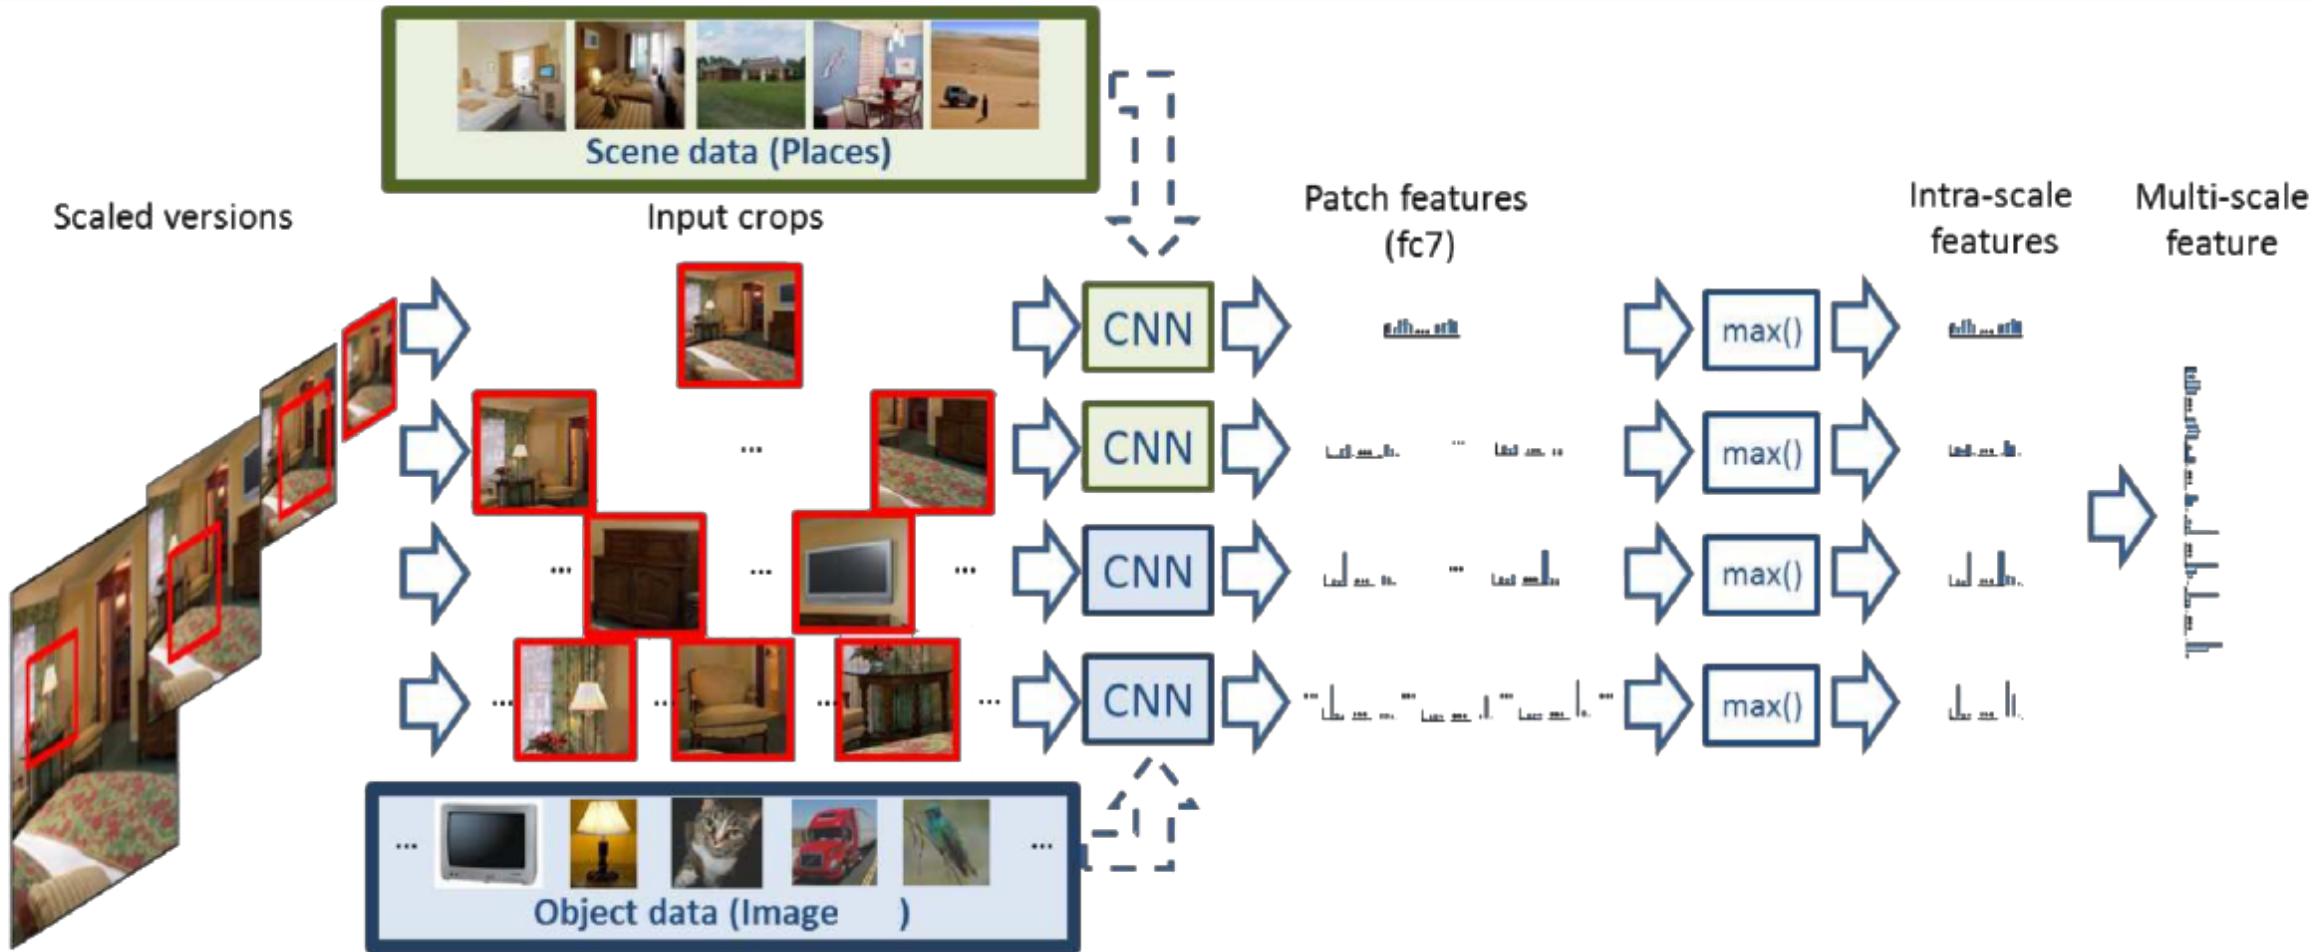
知识点：Fast RCNN , Faster RCNN, Yolo, SSD

#### 第十三课：LSTM 标注学习

#### 第十四课：实践：使用Tensorflow/Keras在数据集上进行目标检测

附加：Kaggle竞赛习题讲解

Python is a must prerequisite



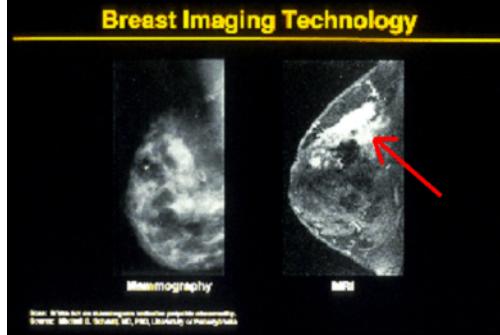
# Table of contents

- General introduction of computer vision
  - The applications of computer vision
  - Definition of computer vision
  - How to learn CV (APWF method)
- Pixel Operations
- Quantilizations
- Wraping/ Combing / Scaling
- Image Convolution

# What Computer Vision matters



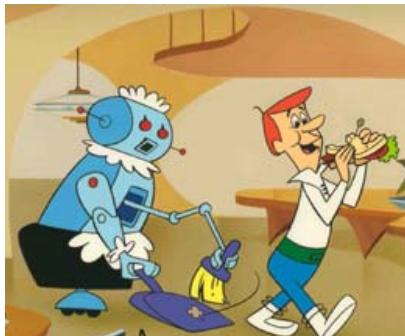
Safety



Health



Security

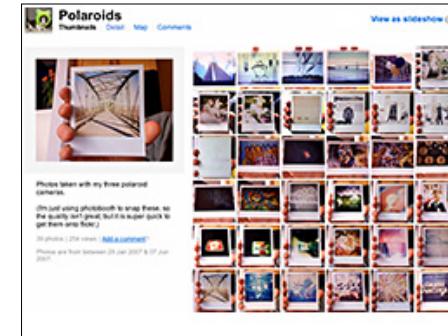


Comfort



Fun

DATAGURU专业数据分析社区



Access

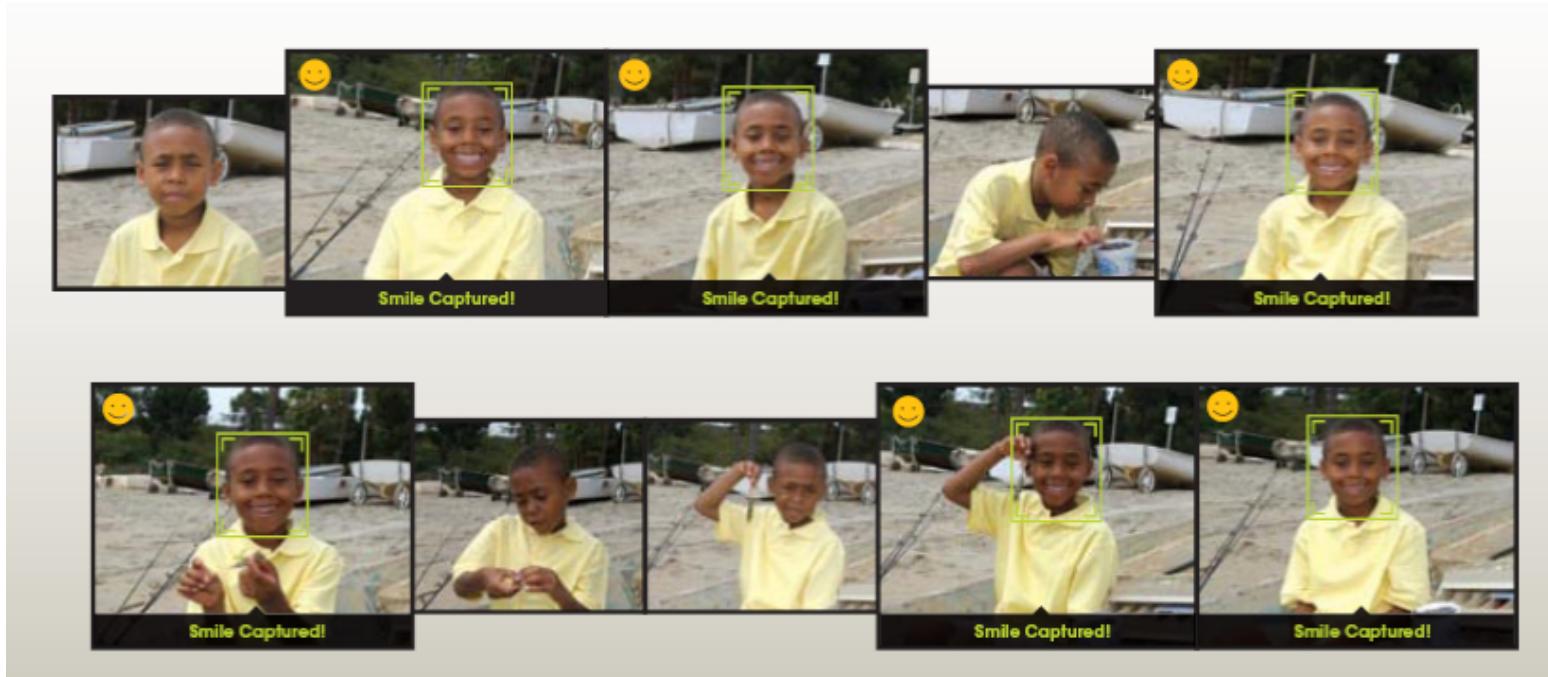
# CV applications

Detection	Recognition	Video Feature Processing	3D Modelling	Robotics	Auto-coding
Smile detection	Object Recognition	Target Recognition	Shape Capture	Medical	GAN
Face detection	Face Verification	Human Feature	Motion Capture	Space	...
	Biometrics	Tracing	Automobile	Sport	
	Picture search	Industry motion	AR/VR		
	OCR	SLAM	3D Reconstruction		

# Smile detection ( FER )

## The Smile Shutter flow

Imagine a camera smart enough to catch every smile! In Smile Shutter Mode, your Cyber-shot® camera can automatically trip the shutter at just the right instant to catch the perfect expression.



## Algorithm Example:

ASM Feature  
SVM or Bayesian Classification  
...

# Face detection



Algorithm Example:

Haar Feature  
Cascade CNN  
DMP Model

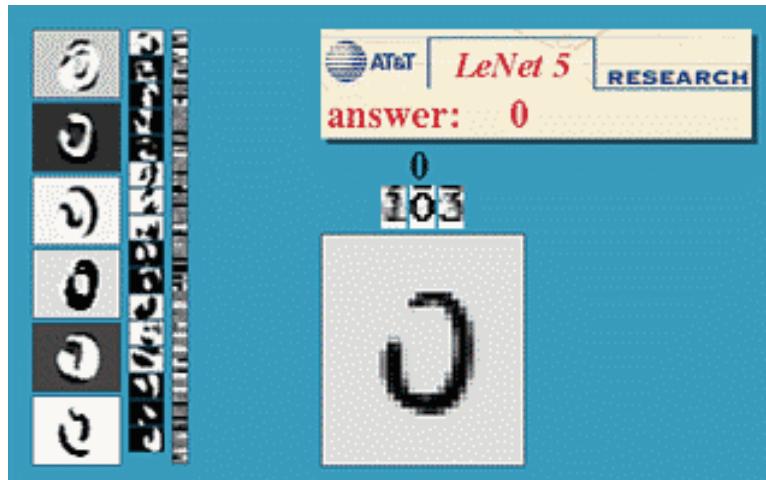
...

- Many new digital cameras and apps now detect faces

## Technology to convert scanned docs to text

- If you have a scanner, it probably came with OCR software

Algorithm Example:



Digit recognition, AT&T labs

<http://www.research.att.com/~yann/>



License plate readers

[http://en.wikipedia.org/wiki/Automatic\\_number\\_plate\\_recognition](http://en.wikipedia.org/wiki/Automatic_number_plate_recognition)

Denoise  
Fast RCNN / Yolo

...

# Object recognition (in supermarkets)



Algorithm Example:

RCNN/Yolo/SSD

...

## LaneHawk by EvolutionRobotics

“A smart camera is flush-mounted in the checkout lane, continuously watching for items. When an item is detected and recognized, the cashier verifies the quantity of items that were found under the basket, and continues to close the transaction. The item can remain under the basket, and with LaneHawk, you are assured to get paid for it...”

DATAGURU专业数据分析社区

# Object recognition (in Smartphone)



Algorithm Example:

RCNN/Yolo/SSD

...

Point & Find, Nokia

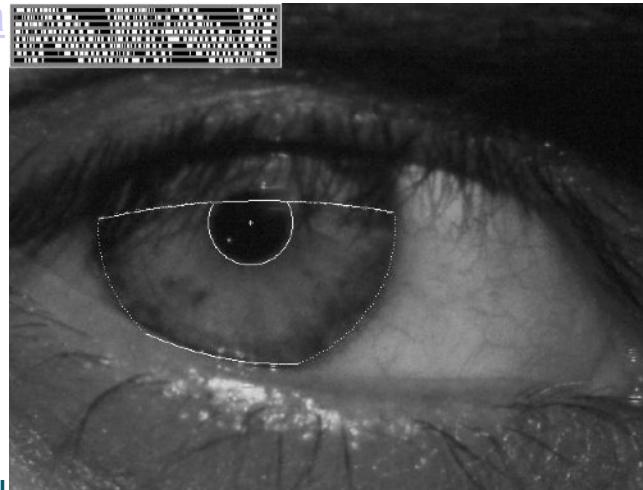
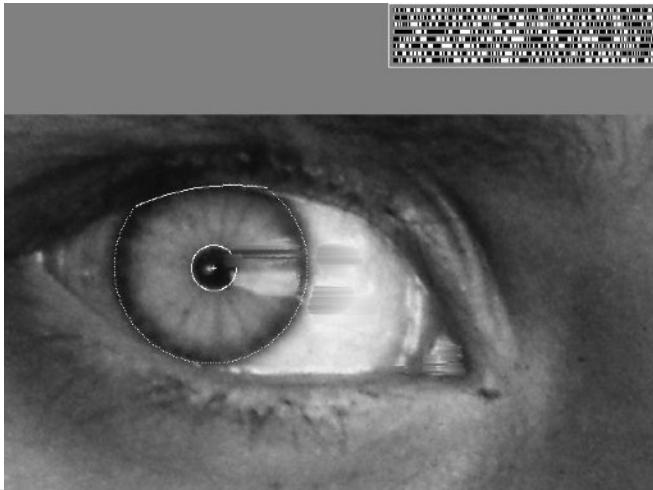
Google Goggles

# Iris Recognition



*“How the Afghan Girl was Identified by Her Iris Patterns”* Read the [story](#)

[wikipedia](#)



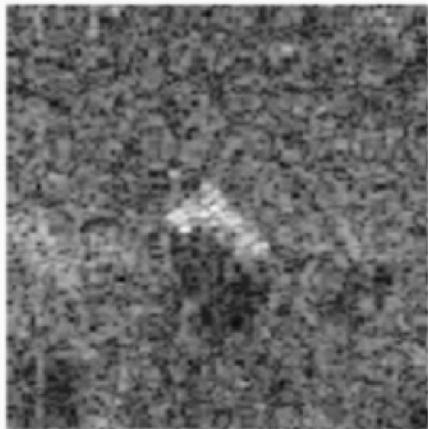
Algorithm Example:

Image Unification

...

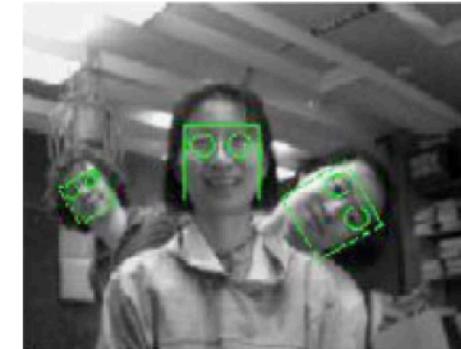
- Target recognition
  - Find enemy vehicles (which are trying not to be found!)
- Human Interfaces
  - Detect faces and identify people
  - Recognize gestures, activities

FLIR image from  
[sdvision.kaist.ac.kr/](http://sdvision.kaist.ac.kr/)



SAR image from  
[web.mit.edu/6.003/courseware/](http://web.mit.edu/6.003/courseware/)

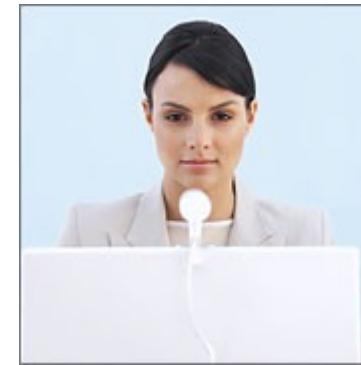
Algorithm Example:  
Lucas-Kanade algorithm  
Skeleton  
...



# Face Verification



Fingerprint scanners on  
many new laptops,  
other devices



Algorithm Example:

Eigenface Feature  
HMM/SVD classification

...

Face recognition systems now  
beginning to appear more widely  
<http://www.sensiblevision.com/>

# Special effects: shape capture



Algorithm Example:

Deepface algorithm

...

*The Matrix* movies, ESC Entertainment, XYZRGB, NRC

DATAGURU专业数据分析社区

# Special effects: motion capture



*Pirates of the Caribbean*, Industrial Light and Magic

Algorithm Example:

Deepface algorithm

...

# Sports



*Sportvision* first down line  
Nice [explanation](#) on [www.howstuffworks.com](http://www.howstuffworks.com)

<http://www.sportvision.com/video.html>

# Intelligence car (unmanned vehicle)

►► manufacturer products      consumer products ◀◀

## Our Vision. Your Safety.



rear looking camera      forward looking camera  
side looking camera

**EyeQ** Vision on a Chip  [read more](#)

**Vision Applications**  Road, Vehicle, Pedestrian Protection and more [read more](#)

**AWS** Advance Warning System  [read more](#)

News

- Mobileye Advanced Technologies Power Volvo Cars World First Collision Warning With Auto Brake System
- Volvo: New Collision Warning with Auto Brake Helps Prevent Rear-end

[all news](#)

Events

- Mobileye at Equip Auto, Paris, France
- Mobileye at SEMA, Las Vegas, NV

[read more](#)

- Mobileye
  - Vision systems currently in high-end BMW, GM, Volvo models

Algorithm Example:

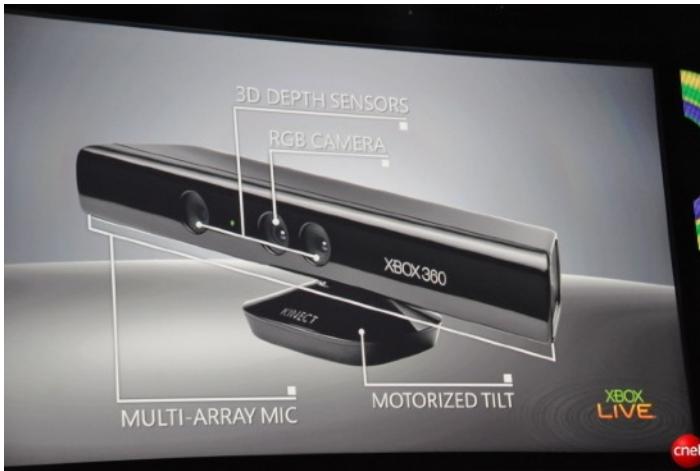
CNN+ SVM decision  
SLAM

...

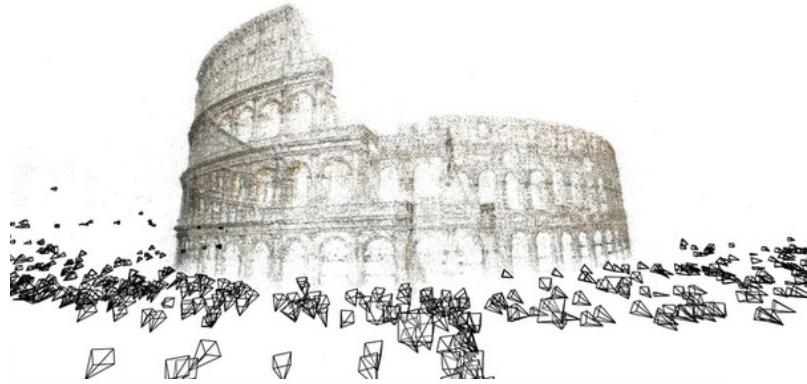
# Intelligence car (unmanned vehicle)



- Object Recognition: <http://www.youtube.com/watch?feature=iv&v=fQ59dXOo63o>
- Mario: <http://www.youtube.com/watch?v=8CTJL5IUjHg>
- 3D: <http://www.youtube.com/watch?v=7QrnwoO1-8A>
- Robot: <http://www.youtube.com/watch?v=w8BmgtMKFbY>



# 3D from thousands of images



Building Rome in a Day: Agarwal et al. 2009



NASA's Mars Exploration Rover Spirit captured this westward view from atop a low plateau where Spirit spent the closing months of 2007.

## Vision systems (JPL) used for several tasks

- Panorama stitching
- 3D terrain modeling
- Obstacle detection, position tracking
- For more, read “Computer Vision on Mars” by Matthies et al.

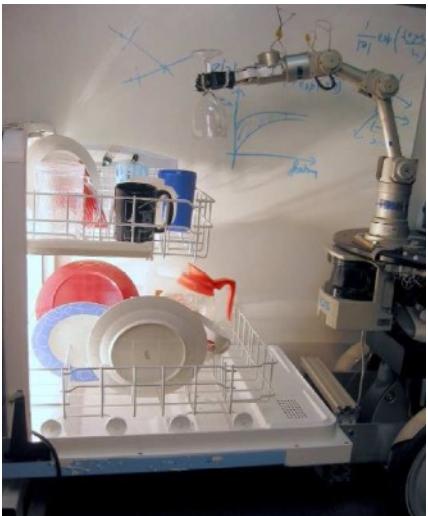
DATAGURU专业数据分析社区

# Industrial robots

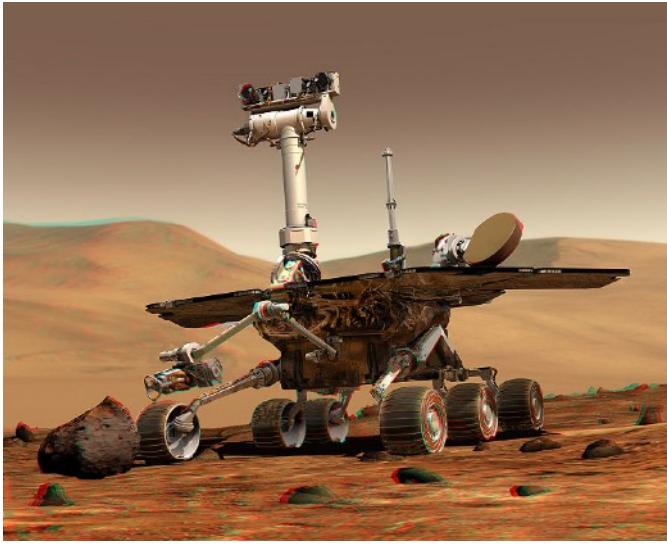


Vision-guided robots position nut runners on wheels

# Mobile robots



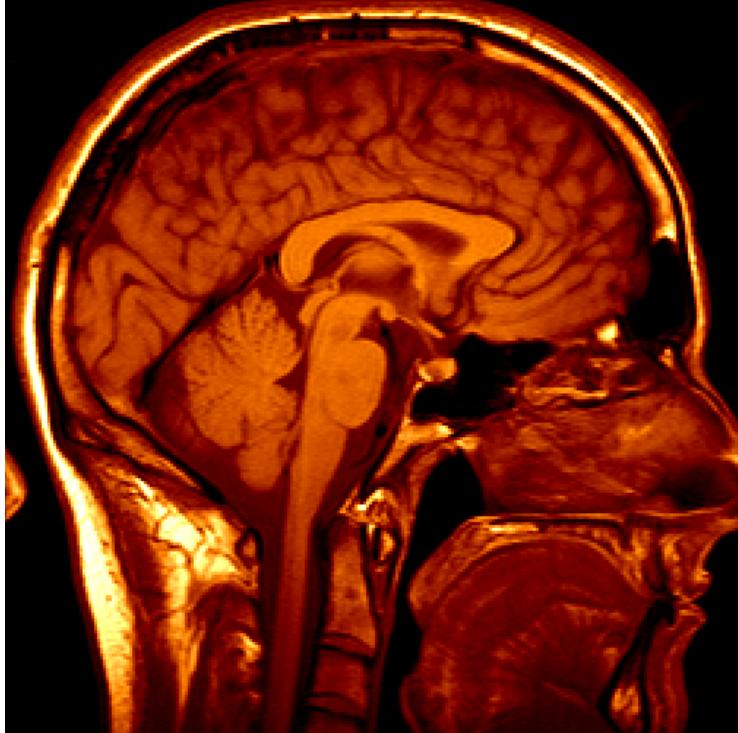
Saxena et al. 2008  
STAIR at Stanford



NASA's Mars Spirit Rover  
[http://en.wikipedia.org/wiki/Spirit\\_rover](http://en.wikipedia.org/wiki/Spirit_rover)



<http://www.robocup.org/>



3D imaging  
MRI, CT



Image guided surgery  
Grimson et al., MIT

# Table of contents

- General introduction of computer vision
  - The applications of computer vision
  - Definition of computer vision
  - How to learn CV (APWF method)
- Pixel Operations
- Quantilizations
- Wraping/ Combing / Scaling
- Image Convolution

# Unsupervised learning



Algorithm Example:

GAN

...

Translate image/video to computer understood semantic



Make computers **understand** what they saw



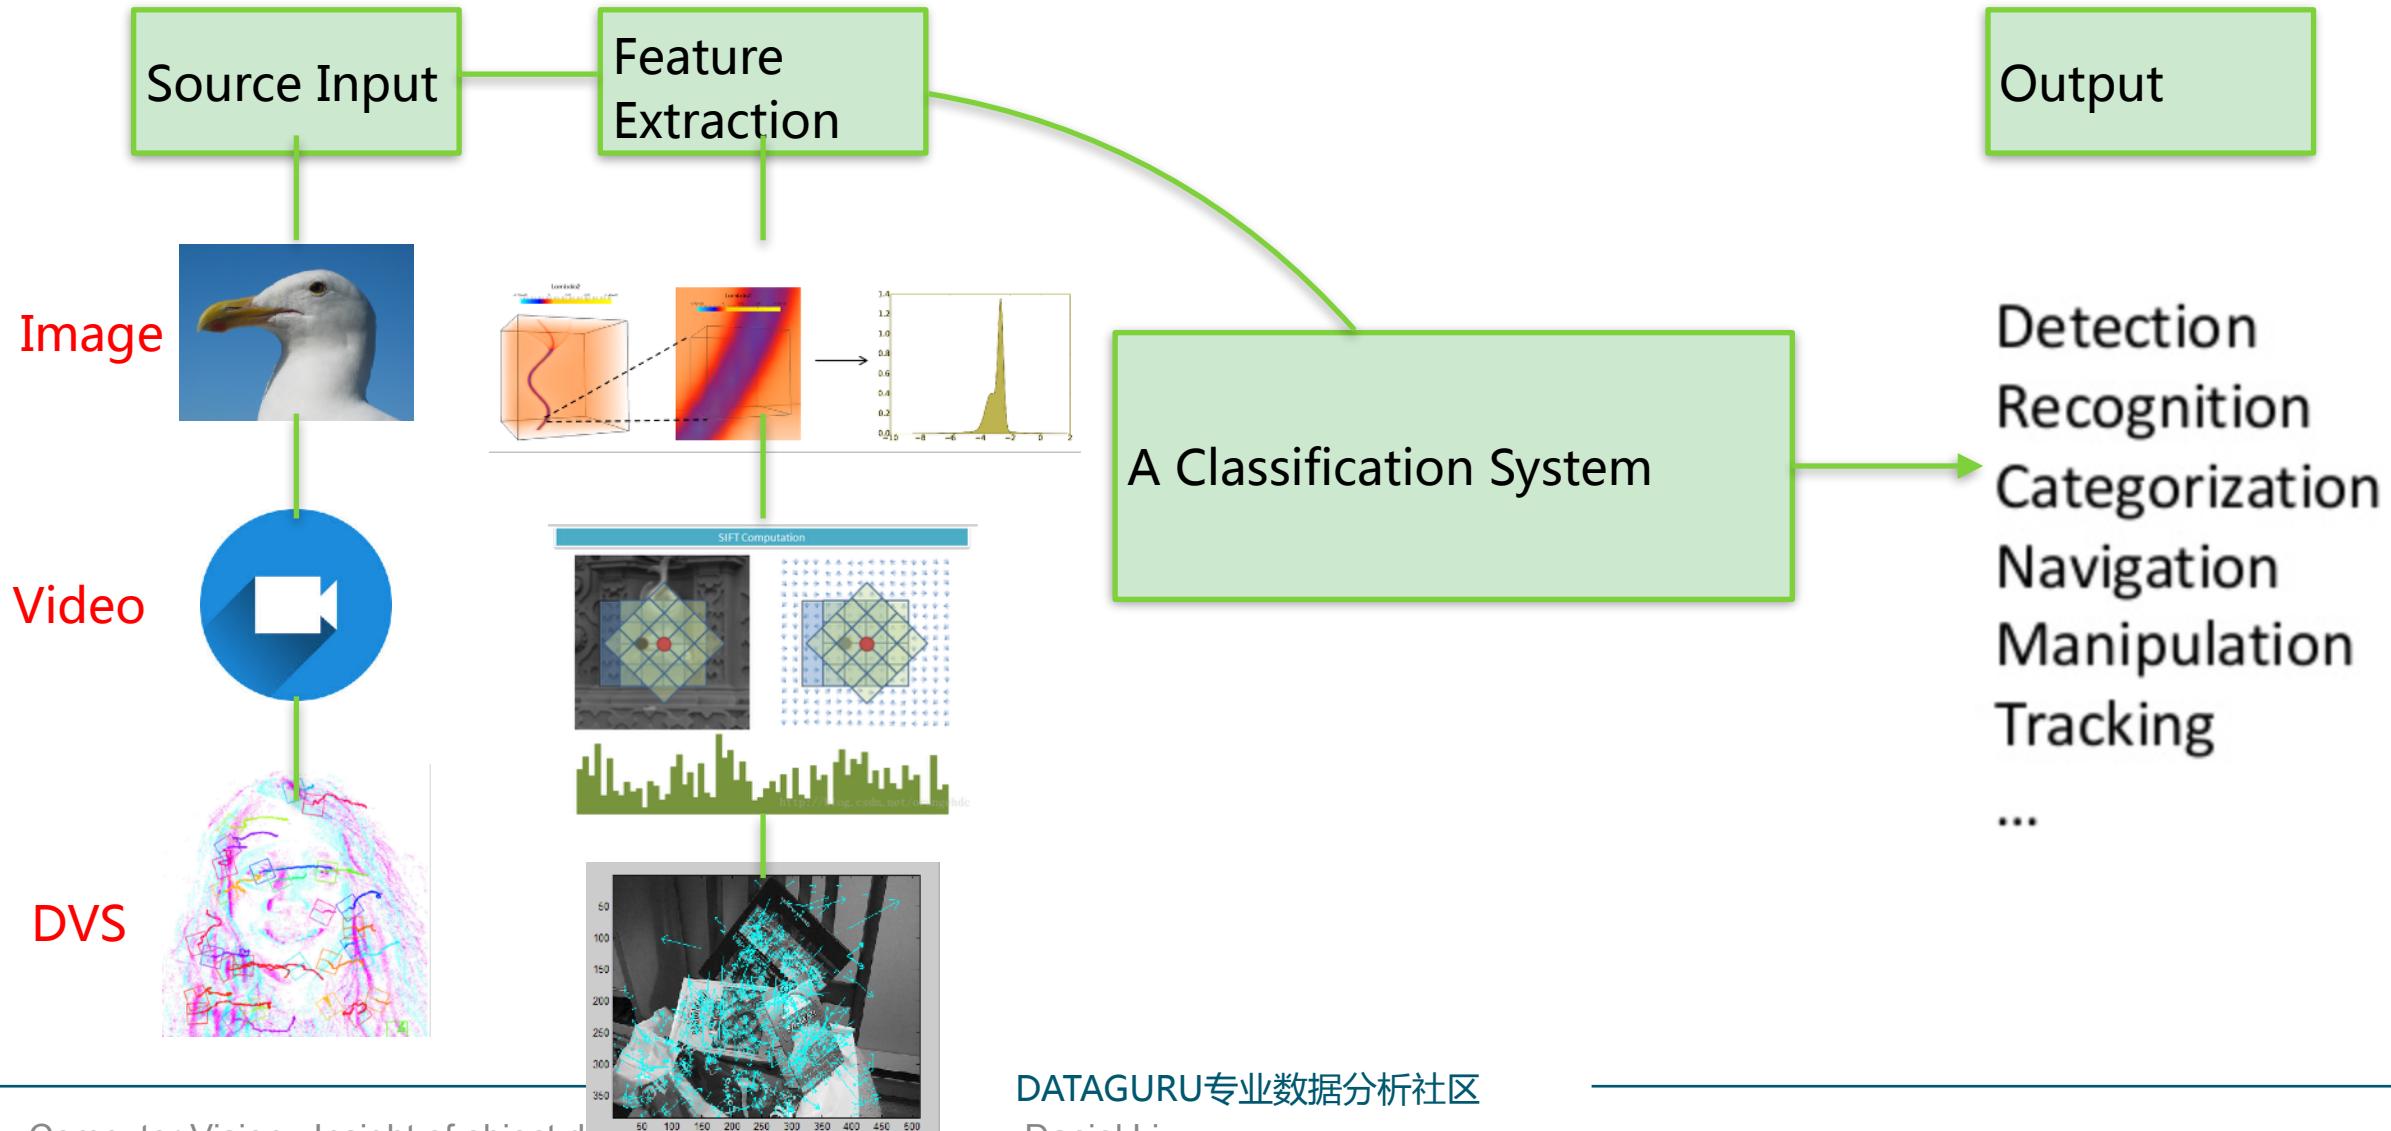
What kind of scene?

Where are the cars?

How far is the building?

...

# A computer vision system



# Table of contents

- General introduction of computer vision
  - The applications of computer vision
  - Definition of computer vision
  - How to learn CV (APWF method)
- Pixel Operations
- Quantilizations
- Wraping/ Combing / Scaling
- Histogram operations

# How to learn CV (AFPW method)



Wheel

```
def DoG_Extrema(Octaves, oc, current):
    top = current - 1
    down = current + 1
    current_state = Octaves[oc][1:-2, 1:-2, current]

    for x in range(0, 3):
        for y in range(0, 3):
            if not (x == 1 and y == 1):
                local_maxima_to_current_neighbors = (current_state > Octaves[oc][x:x - 3, y:y - 3, current])
                local_maxima_to_top_neighbors = (current_state > Octaves[oc][x:x - 3, y:y - 3, top])
                local_maxima_to_down_neighbors = (current_state > Octaves[oc][x:x - 3, y:y - 3, down])

            try:
                local_maxima = local_maxima & local_maxima_to_current_neighbors & local_maxima_to_top_neighbors & local_maxima_to_down_neighbors
            except:
                local_maxima = local_maxima_to_current_neighbors

    for x in range(0, 3):
        for y in range(0, 3):
            if not (x == 1 and y == 1):
                local_minima_to_current_neighbors = (current_state < Octaves[oc][x:x - 3, y:y - 3, current])
                local_minima_to_top_neighbors = (current_state < Octaves[oc][x:x - 3, y:y - 3, top])
                local_minima_to_down_neighbors = (current_state < Octaves[oc][x:x - 3, y:y - 3, down])
            try:
                local_minima = local_minima & local_minima_to_current_neighbors & local_minima_to_top_neighbors & local_minima_to_down_neighbors
            except:
                local_minima = local_minima_to_current_neighbors
```

Perception

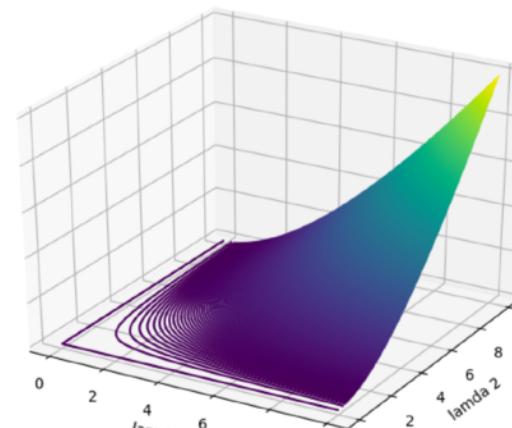
$$\frac{\det(M)}{\text{trace}(M)} = \frac{I_x^2 I_y^2 - (I_x I_y)^2}{I_x^2 + I_y^2}$$

Algorithms

```
im = cv2.imread('Images/jobs.jpg')
SIFT = cv2.xfeatures2d.SIFT_create()
SURF = cv2.xfeatures2d.SURF_create()
ORB = cv2.ORB_create()
```

```
kps_SIFT,descfs_SIFT = SIFT.detectAndCompute(im, None)
kps_SURF,descfs_SURF = SURF.detectAndCompute(im, None)
kps_ORB, descfs_ORB = ORB.detectAndCompute(im, None)
```

```
imgSIFT = cv2.drawKeypoints(im, kps_SIFT, None, color=(255,0,0))
imgSURF = cv2.drawKeypoints(imgSIFT, kps_SURF, None, color=(0,255,0))
imgORB = cv2.drawKeypoints(imgSURF, kps_ORB, None, color=(0,0,255))
```



DAIGURU专业数据分析社区

Functions

# AFPW Method Example:

Wheel

```
def sigmoid(x):
    sigmoid = 1 / (1 + np.exp(-x))
    return sigmoid
```

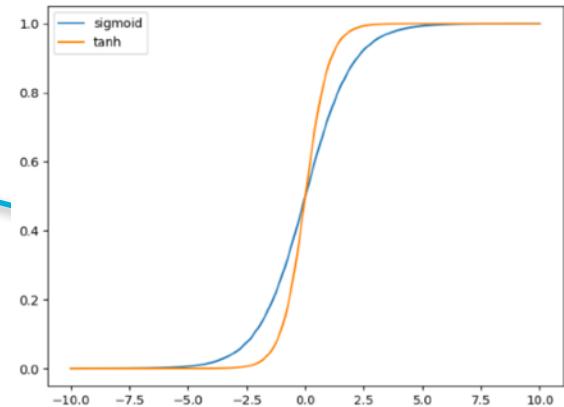
Perception

$$\text{sigmoid} = \frac{1}{1 + e^{-\lambda z}}$$

Algorithms

tensorflow.sigmoid

Functions



# Table of contents

- General introduction of computer vision
- Pixel Operations
  - Pixel Operation overview
  - Pixel changing
- Quantilizations
- Wrapping/ Combing / Scaling
- Histogram operations

# Overview



Pixel Operation

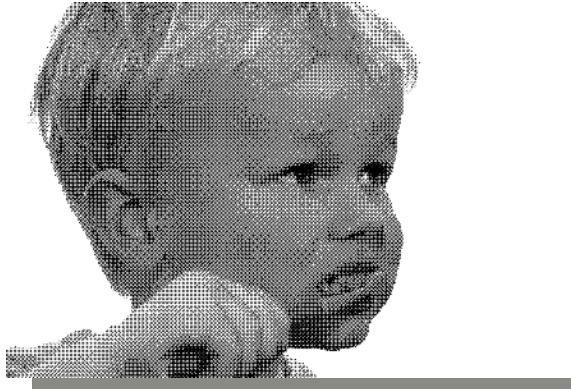


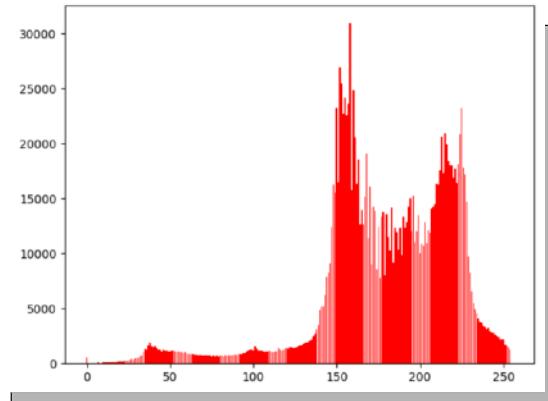
Image Quantilization



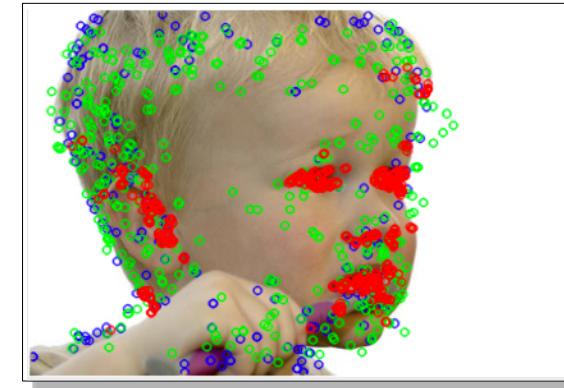
Wrap/Affine/combining operations



Image Convolution



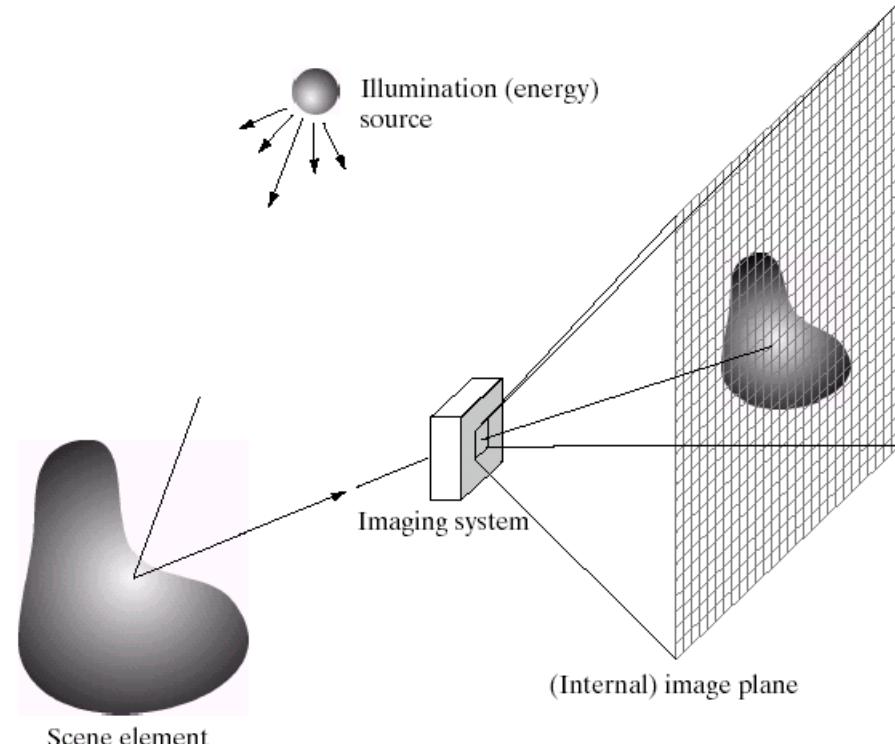
Histogram Operation



Interesting point Extraction

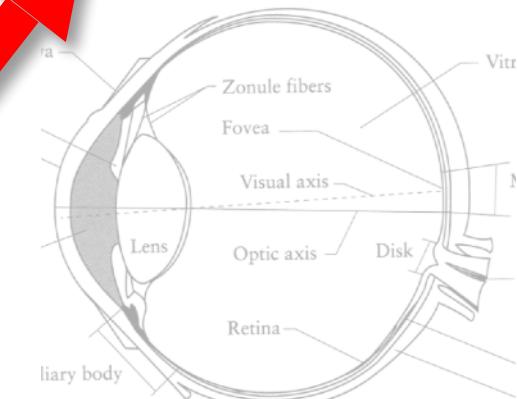
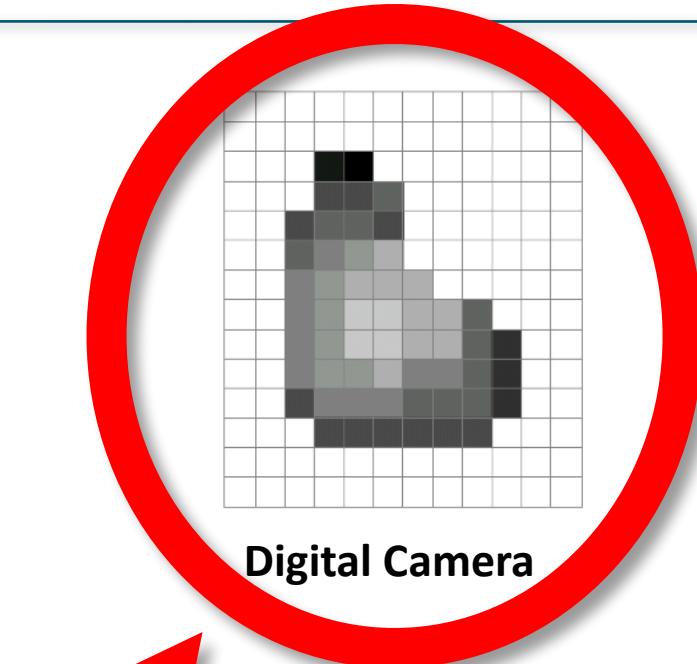
- Pixel operations
  - Add random noise
  - Add luminance
  - Add contrast
  - Add saturation
- Quantization
  - Uniform quantization
  - Random dither
  - Ordered dither
  - Floyd-Steinberg dither
- Warping/Combining/Warping
  - Scale
  - Rotate
  - Warps
- Convolution
  - Blur
  - Detect edge
  - Haar Feature
- Histogram
  - Histogram equalisation
  - Hog Feature
- Interesting points
  - Haris corner point
  - SIFT/ORB/SURF
  - Descriptor match

# What is an image?



We'll focus on these in this class

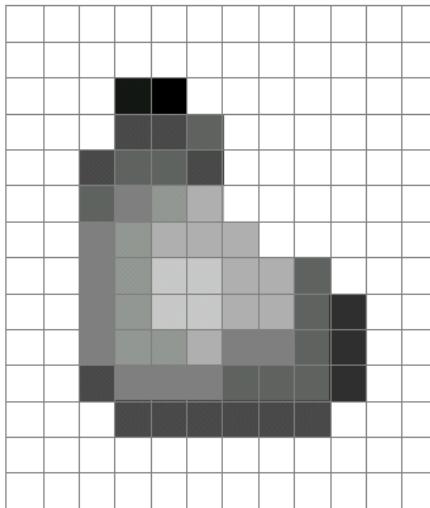
(More on this process later)



The Eye

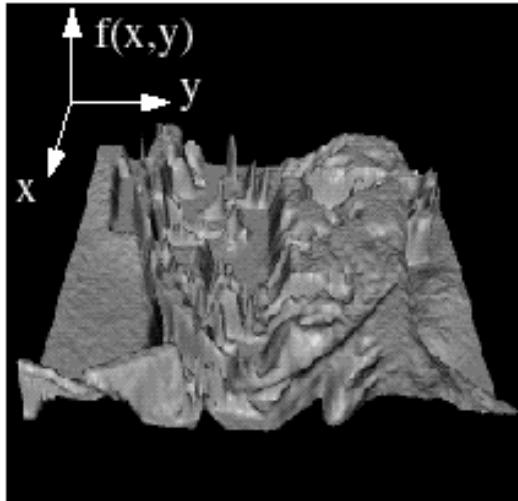
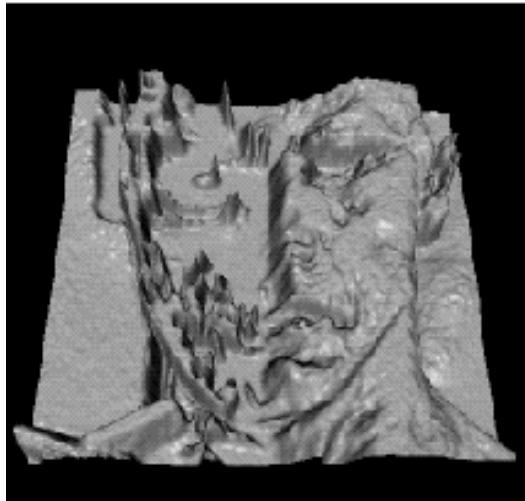
Source: A. Efros

# What is an image?



255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255	255	255	255	255
255	255	127	145	145	175	127	127	95	47	255	255	255	255	255	255
255	255	74	127	127	127	95	95	95	47	255	255	255	255	255	255
255	255	255	74	74	74	74	74	74	74	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255

# Image as a Function



- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$



$$g(x,y) = f(-x,y)$$

- We'll talk about a special kind of operator, *convolution* (linear filtering)

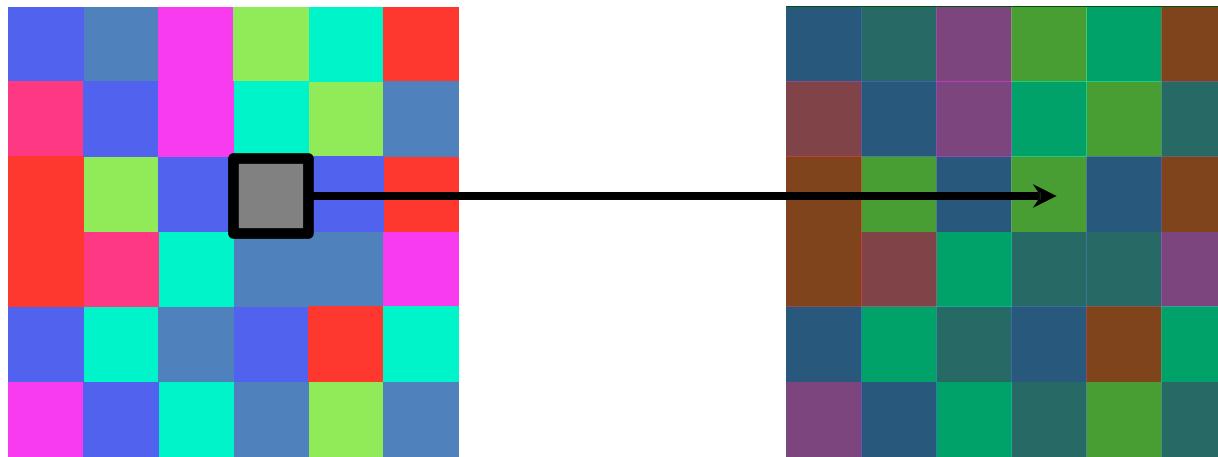
# Image as a Function

- We can think of an **image** as a function,  $f$ ,
- $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ 
  - $f(x, y)$  gives the **intensity** at position  $(x, y)$
  - Realistically, we expect the image only to be defined over a rectangle, with a finite range:
    - $f: [a,b] \times [c,d] \rightarrow [0,1]$
- A color image is just three functions pasted together. We can write this as a “vector-valued” function:

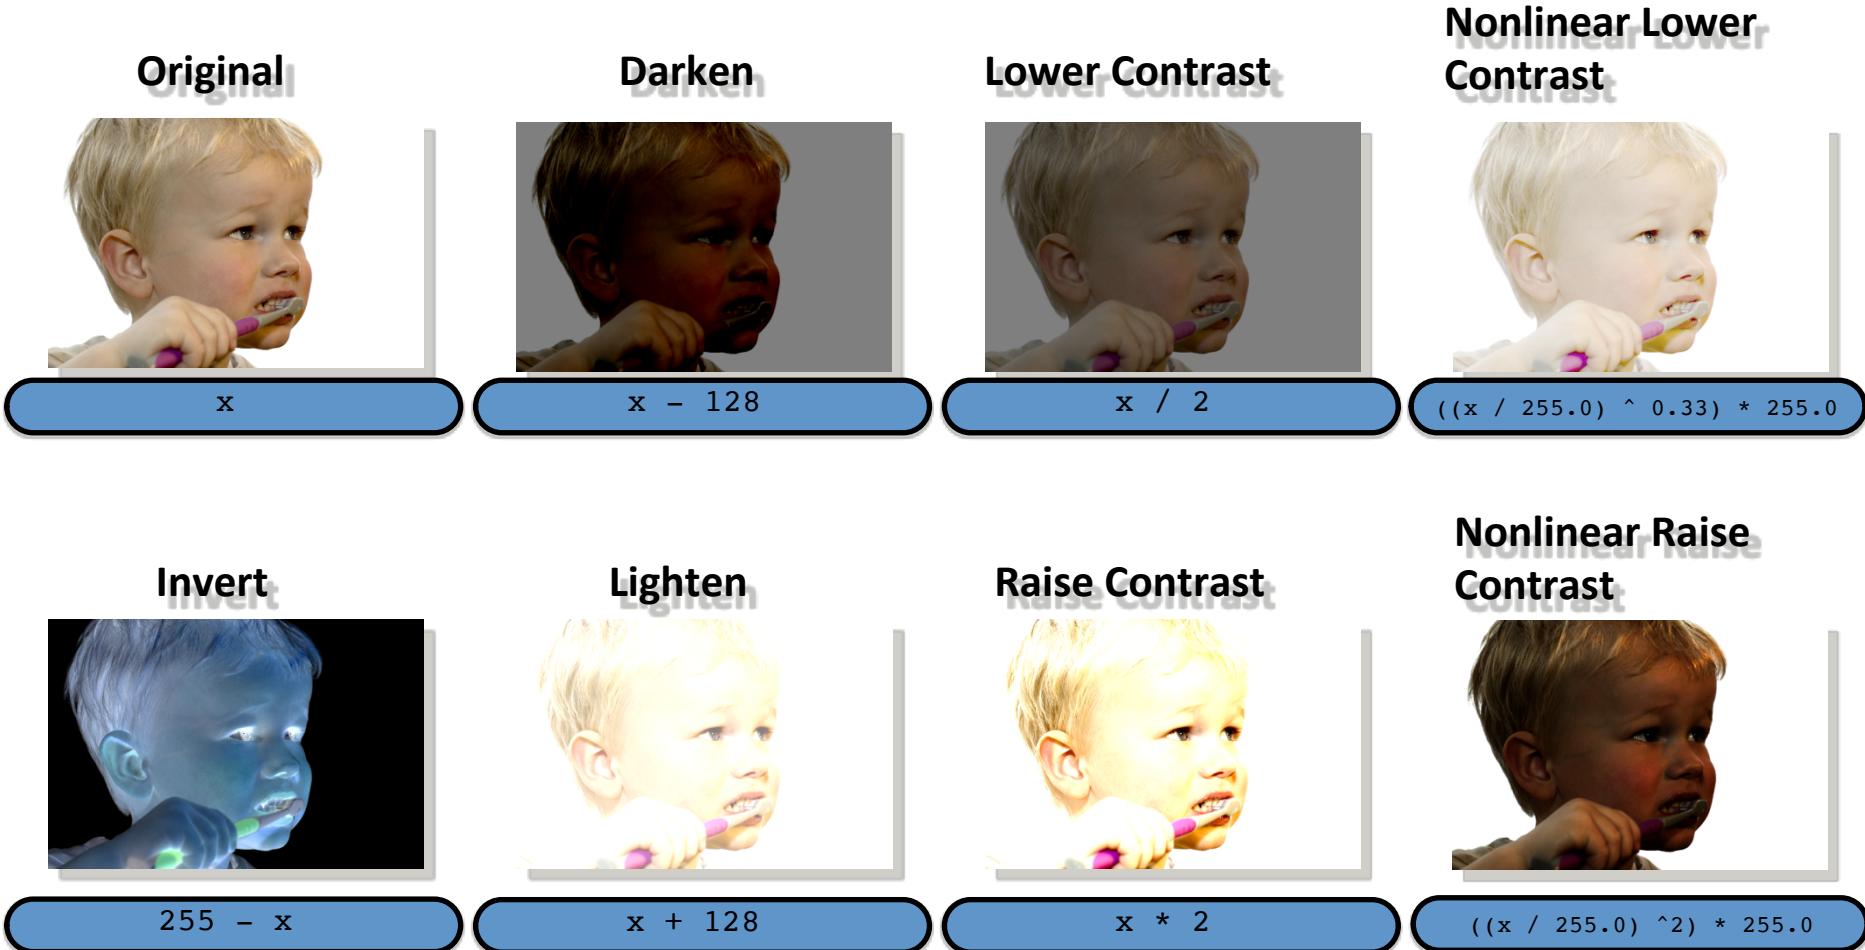
$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

DATAGURU专业数据分析社区

# Point Operations



# Point Processing



- Some operations preserve the range but change the domain of  $f$ :

$$g(x, y) = f(t_x(x, y), t_y(x, y))$$

What kinds of operations can this perform?

- Still other operations operate on both the domain and the range of  $f$ .

# Table of contents

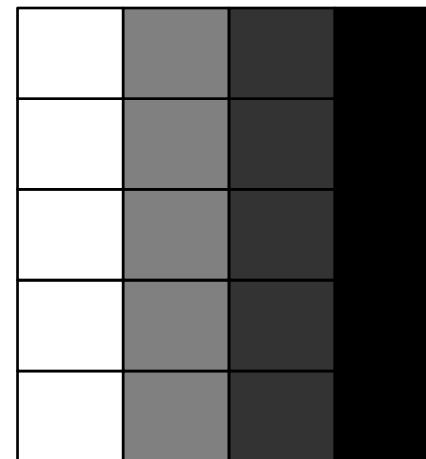
- General introduction of computer vision
- Pixel Operations
- Quantilizations
  - Uniform Quantilization
  - Ordered Dithering
  - Halftoning
  - Floyd Steinberg Dithering
- Wrapping/ Combing / Scaling
- Image Convolution

# Quantization

Artifact due to limited intensity resolution

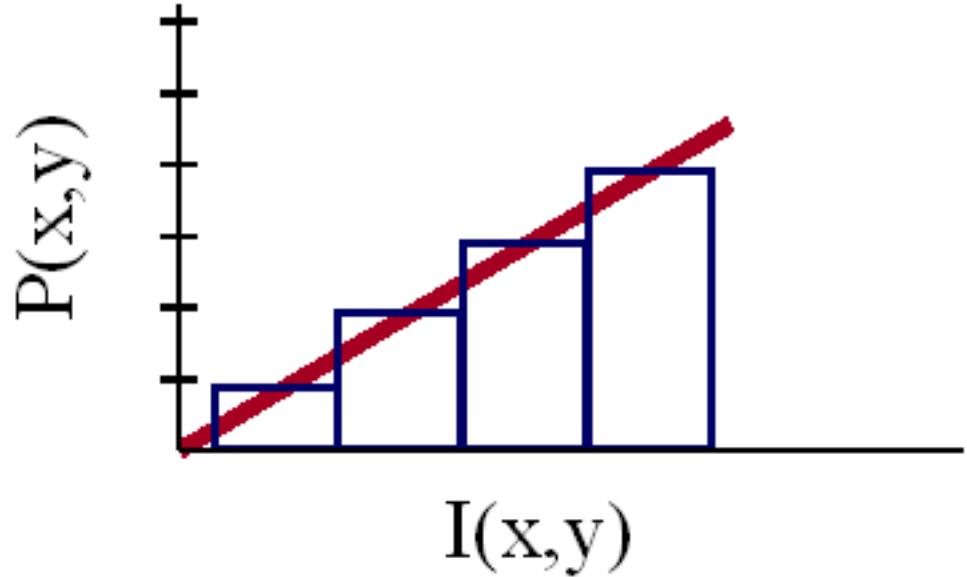
Frame buffers have limited number of bits per pixel

Physical devices have limited dynamic range



# Uniform Quantization

$$P(x, y) = \text{trunc}(I(x, y) + 0.5)$$



$I(x, y)$



$P(x, y)$   
2 bits per pixel

# Uniform Quantization

Image with decreasing bits per pixel:



8 bits



4 bits



2 bits



1 bit

Notice contouring

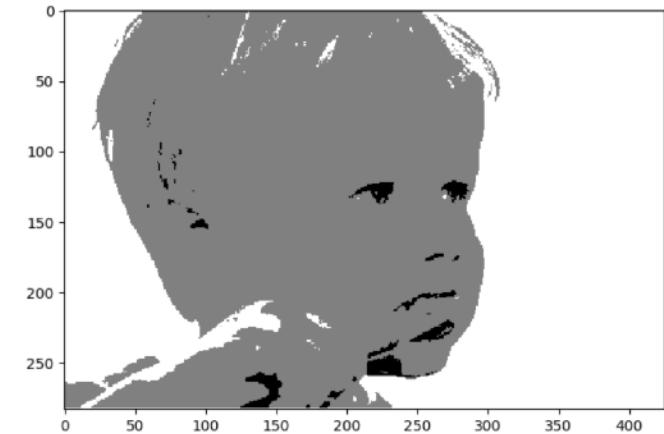
# Naïve Threshold Algorithm



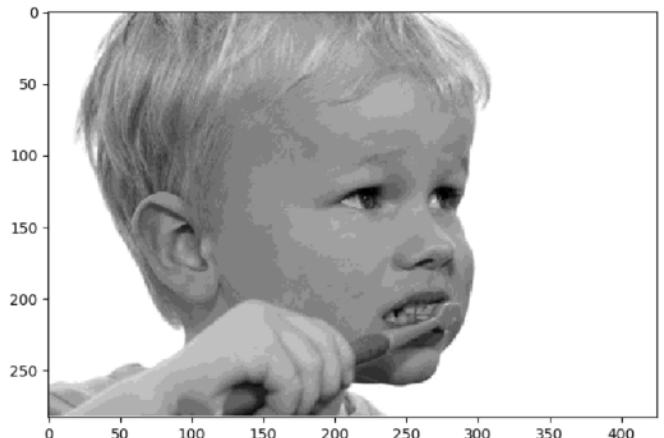
# Coding Example

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 N = 2
6 gray = cv2.imread('images/ChildBrushingTeeth.jpg',cv2.IMREAD_GRAYSCALE)
7 Uniform_quant = np.round(gray*(N/255))*np.round(255/N)
8 print(Uniform_quant)
9
10 plt.imshow(Uniform_quant,cmap='gray')
11 plt.show()
```

N=2



N=16



# Table of contents

- General introduction of computer vision
- Pixel Operations
- Quantilizations
  - Uniform Quantilization
  - Ordered Dithering
  - Halftoning
  - Floyd Steinberg Dithering
- Wrapping/ Combing / Scaling
- Histogram operations

# Dithering

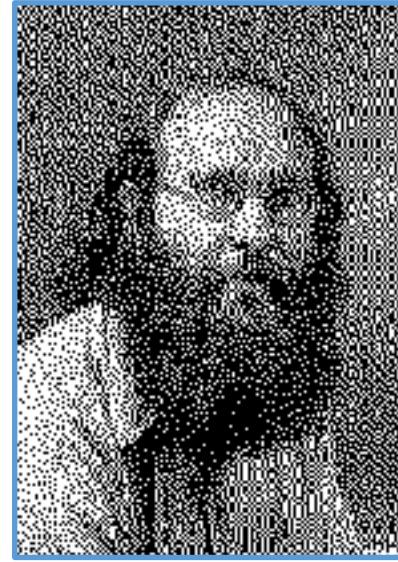
- Distribute errors among pixels
  - Exploit spatial integration in our eye
  - Display greater range of perceptible intensities



Original  
(8 bits)



Uniform  
Quantization  
(1 bit)



Floyd-Steinberg  
Dither  
(1 bit)

# Ordered Dither



Original  
(8 bits)



Uniform  
Quantization  
(1 bit)



Ordered  
Dither  
(1 bit)

# Ordered Dither

ordered dither matrices

$$D_n = \begin{bmatrix} 4D_{\frac{n}{2}} + D_2(1,1)U_{\frac{n}{2}} & 4D_{\frac{n}{2}} + D_2(1,2)U_{\frac{n}{2}} \\ 4D_{\frac{n}{2}} + D_2(2,1)U_{\frac{n}{2}} & 4D_{\frac{n}{2}} + D_2(2,2)U_{\frac{n}{2}} \end{bmatrix}$$

$$D_2 = \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix} \quad D_4 = \begin{bmatrix} 15 & 7 & 13 & 5 \\ 3 & 11 & 1 & 9 \\ 12 & 4 & 14 & 6 \\ 0 & 8 & 2 & 10 \end{bmatrix}$$

# Ordered Dithering

- Break the image into small blocks
- Define a *threshold matrix*
  - Use a different threshold for each pixel of the block
  - Compare each pixel to its own threshold
- The thresholds can be clustered, which looks like newsprint
- The thresholds can be “random” which looks better

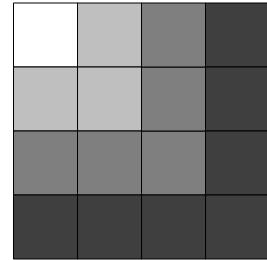
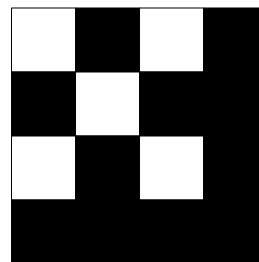


Image block

$$\begin{bmatrix} 1 & 0.75 & 0.5 & 0.25 \\ 0.75 & 0.75 & 0.5 & 0.25 \\ 0.5 & 0.5 & 0.5 & 0.25 \\ 0.25 & 0.25 & 0.25 & 0.25 \end{bmatrix}$$

Threshold matrix

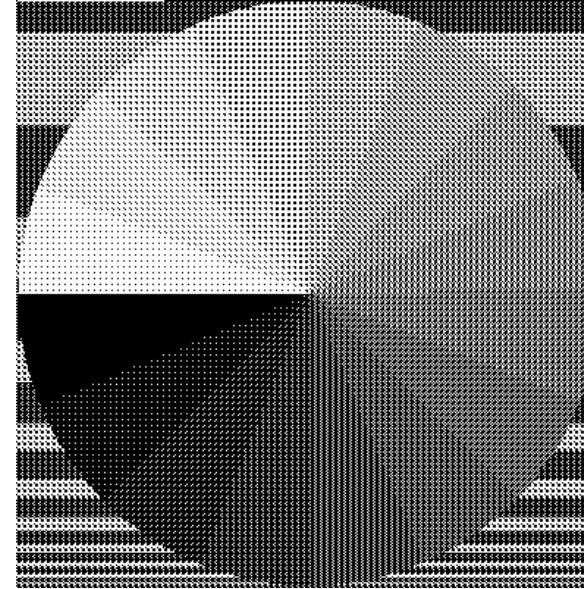
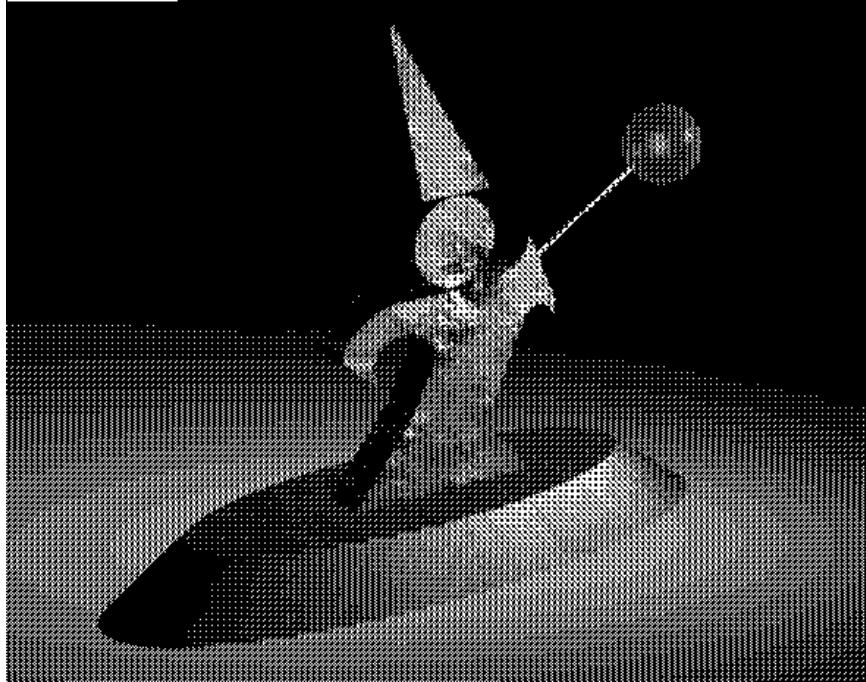
$$\frac{1}{16} \begin{bmatrix} 2 & 16 & 3 & 13 \\ 10 & 6 & 11 & 7 \\ 4 & 14 & 1 & 15 \\ 12 & 8 & 9 & 5 \end{bmatrix}$$



Result

$$\begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

# Ordered Dither



$$\begin{bmatrix} .75 & .375 & .625 & .25 \\ .0625 & 1 & .875 & .4375 \\ .5 & .8125 & .9375 & .125 \\ .1875 & .5625 & .3125 & .6875 \end{bmatrix}$$

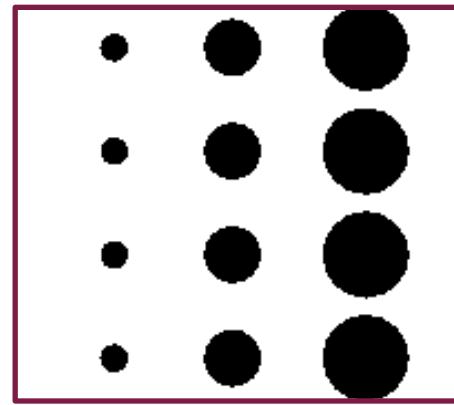
# Table of contents

- General introduction of computer vision
- Pixel Operations
- Quantilizations
  - Uniform Quantilization
  - Ordered Dithering
  - Halftoning
  - Floyd Steinberg Dithering
- Wrapping/ Combing / Scaling
- Image Convolution

- Use dots of varying size to represent intensities
  - Area of dots proportional to intensity in image



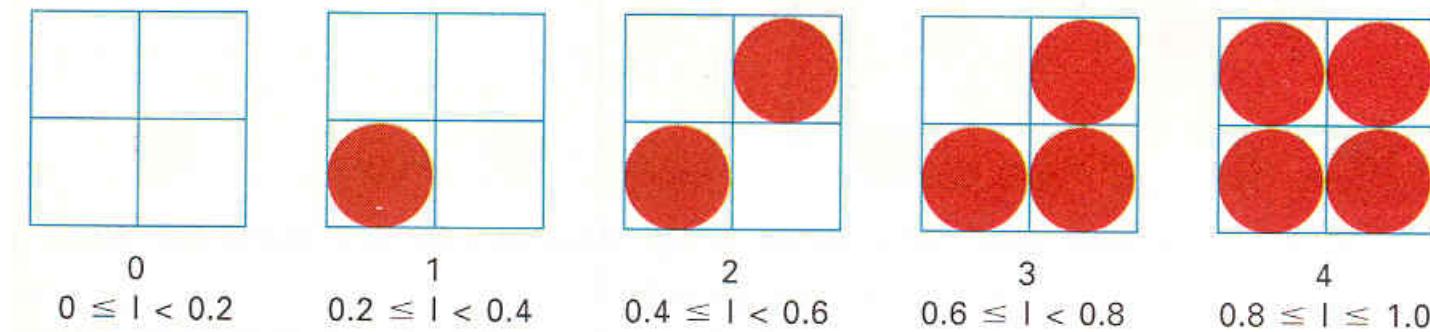
$I(x, y)$



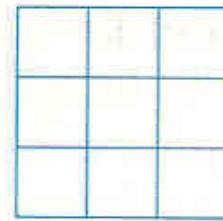
$P(x, y)$

# Classical Halftoning

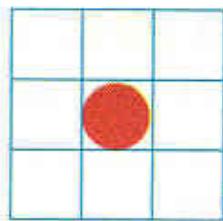
- Use cluster of pixels to represent intensity
  - Trade spatial resolution for intensity resolution



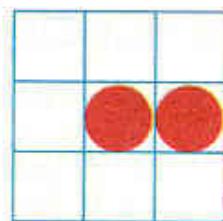
- How many intensities in a  $n \times n$  cluster?



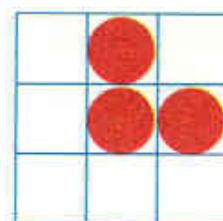
0  
 $0 \leq I < 0.1$



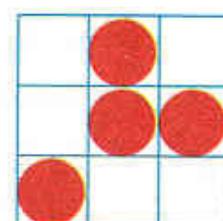
1  
 $0.1 \leq I < 0.2$



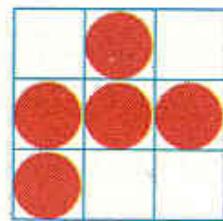
2  
 $0.2 \leq I < 0.3$



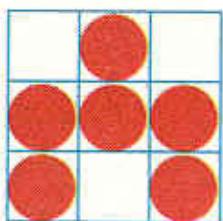
3  
 $0.3 \leq I < 0.4$



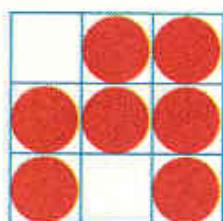
4  
 $0.4 \leq I < 0.5$



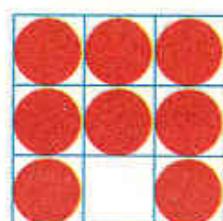
5  
 $0.5 \leq I < 0.6$



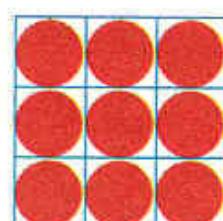
6  
 $0.6 \leq I < 0.7$



7  
 $0.7 \leq I < 0.8$



8  
 $0.8 \leq I < 0.9$

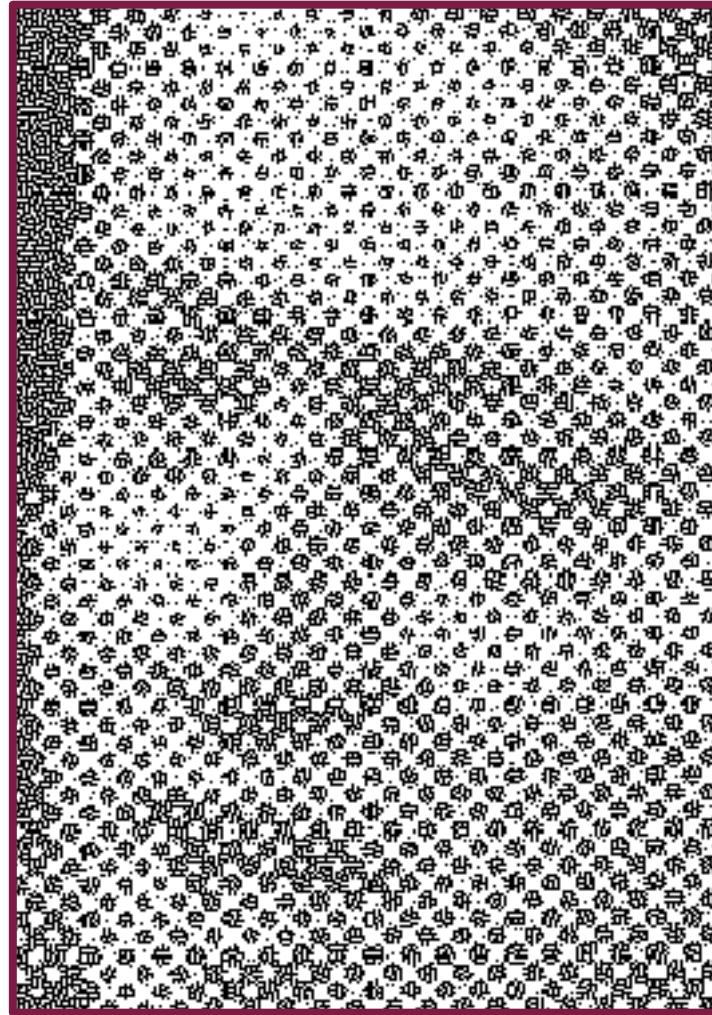


9  
 $0.9 \leq I \leq 1.0$

# Classical Halftoning



Newspaper image



# Table of contents

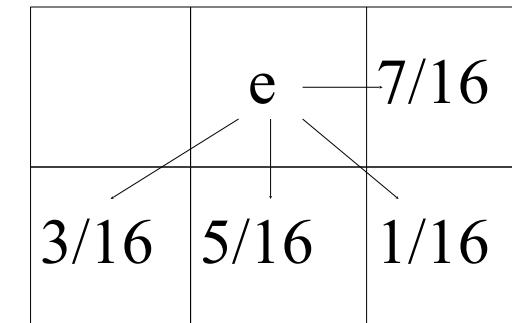
- General introduction of computer vision
- Pixel Operations
- Quantilizations
  - Uniform Quantilization
  - Ordered Dithering
  - Halftoning
  - Floyd Steinberg Dithering
- Wrapping/ Combing / Scaling
- Histogram operations

- Any “rounding” errors are distributed to other pixels
  - Specifically to the pixels below and to the right
    - 7/16 of the error to the pixel to the right
    - 3/16 of the error to the pixel to the lower left
    - 5/16 of the error to the pixel below
    - 1/16 of the error to the pixel to the lower right
- Assume the 1 in the middle gets “rounded” to 0

$$\begin{bmatrix} 0.00 & 0.00 & 0.00 \\ 0.00 & 1.00 & 0.00 \\ 0.00 & 0.00 & 0.00 \end{bmatrix} \longrightarrow \begin{bmatrix} 0.00 & 0.00 & 0.00 \\ 0.00 & 0.00 & 0.44 \\ 0.19 & 0.31 & 0.06 \end{bmatrix}$$

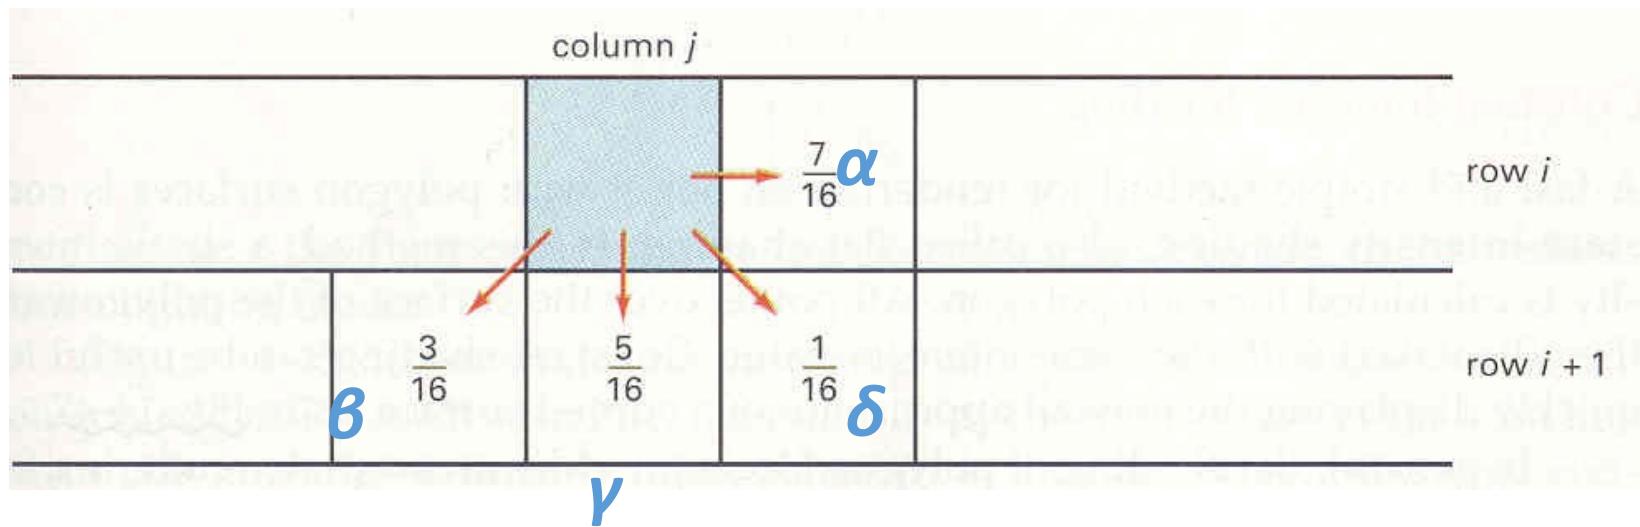
# Floyd-Steinberg Dithering

- Start at one corner and work through image pixel by pixel
  - Usually scan top to bottom in a zig-zag
- Threshold each pixel
- Compute the error at that pixel: The difference between what should be there and what you did put there
  - If you made the pixel 0,  $e = \text{original}$ ; if you made it 1,  $e = \text{original}-1$
- Propagate error to neighbors by adding some proportion of the error to each unprocessed neighbor
  - A *mask* tells you how to distribute the error
- Easiest to work with floating point image
  - Convert all pixels to 0-1 floating point



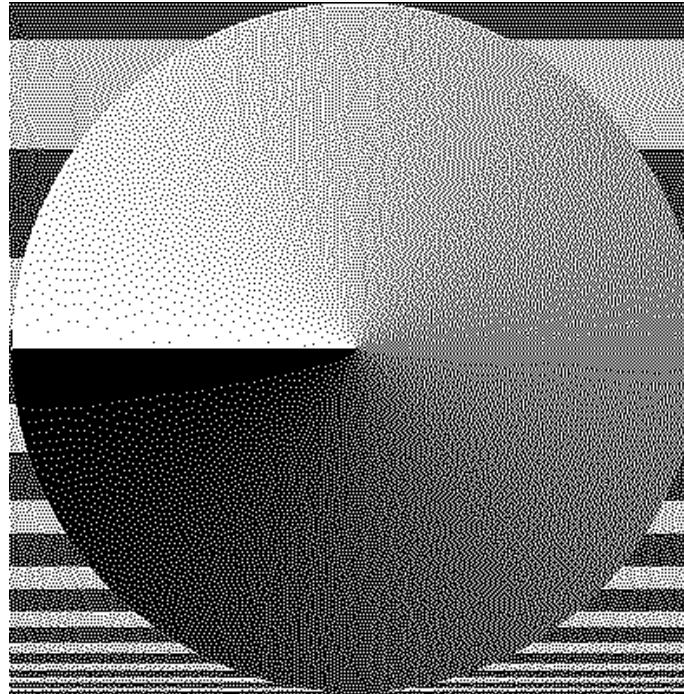
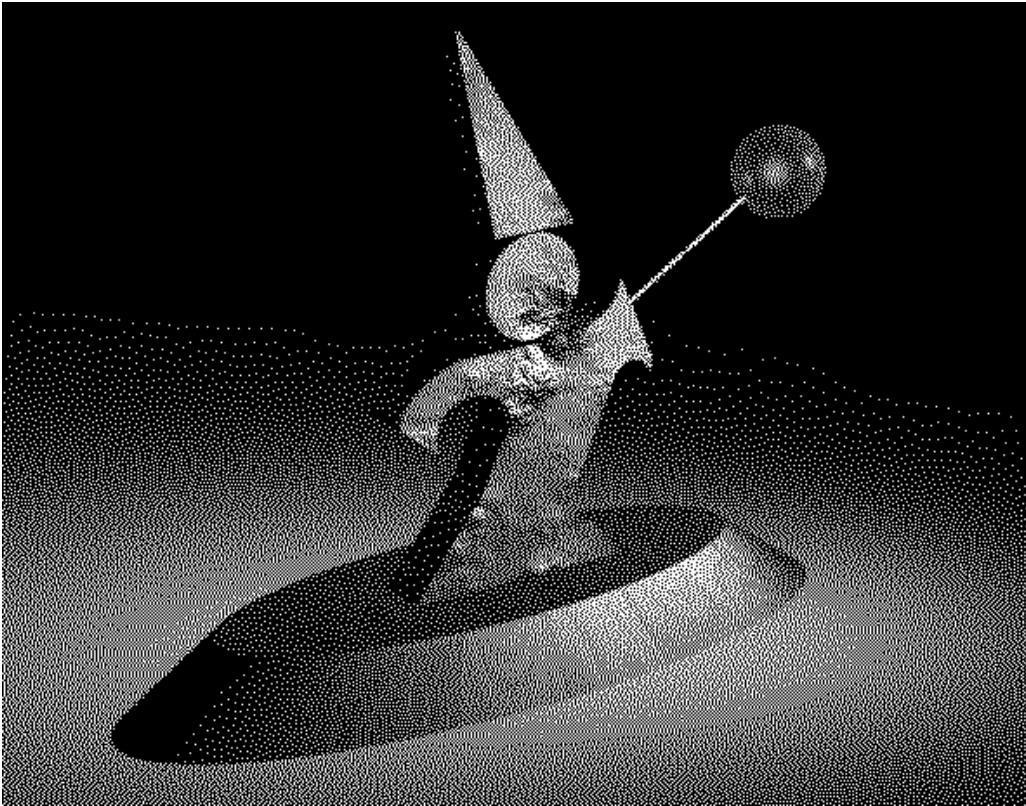
# Error Diffusion Dither

- Spread quantization error over neighbor pixels
  - Error dispersed to pixels right and below



$$\alpha + \beta + \gamma + \delta = 1.0$$

# Floyd-Steinberg Dithering



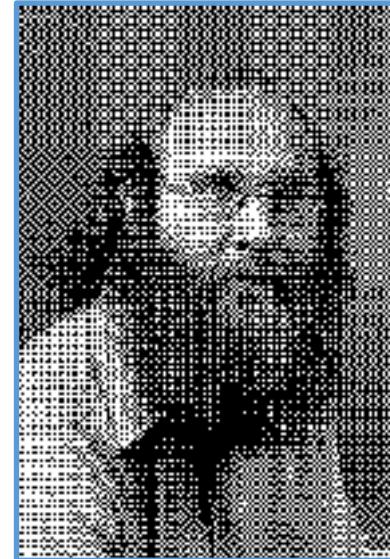
# Error Diffusion Dither



Original  
(8 bits)



Random  
Dither  
(1 bit)



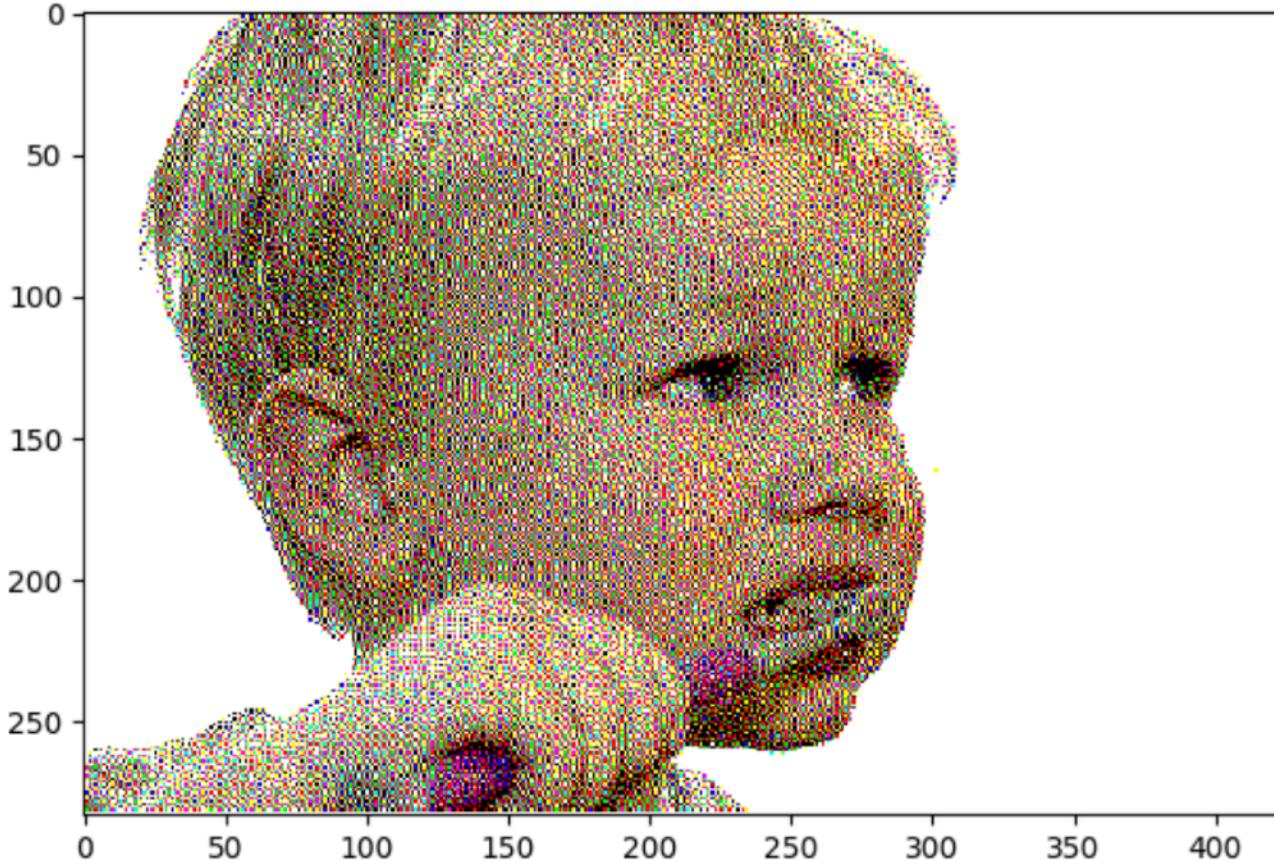
Ordered  
Dither  
(1 bit)



Floyd-Steinberg  
Dither  
(1 bit)

# Dithering Code

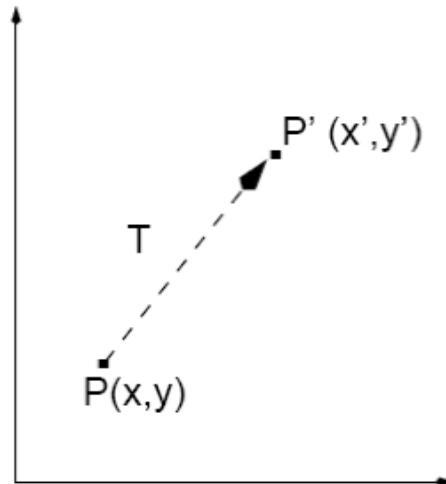
```
1 import numpy as np
2 import PIL.Image as Image
3 import matplotlib.pyplot as plt
4
5 def floyd_steinberg_dither(image_file):
6     new_img = Image.open(image_file)
7     new_img = new_img.convert('RGB')
8     pixel = new_img.load()
9     x_lim, y_lim = new_img.size
10    for y in range(1, y_lim):
11        for x in range(1, x_lim):
12            red_oldpixel, green_oldpixel, blue_oldpixel = pixel[x, y]
13            red_newpixel = np.int(255 * np.floor(red_oldpixel / 128))
14            green_newpixel = np.int(255 * np.floor(green_oldpixel / 128))
15            blue_newpixel = np.int(255 * np.floor(blue_oldpixel / 128))
16
17            pixel[x, y] = red_newpixel, green_newpixel, blue_newpixel
18
19            red_error = red_oldpixel - red_newpixel
20            blue_error = blue_oldpixel - blue_newpixel
21            green_error = green_oldpixel - green_newpixel
22            print(red_error)
23            if x < x_lim - 1:
24                red = pixel[x + 1, y][0] + round(red_error * 7 / 16)
25                green = pixel[x + 1, y][1] + round(green_error * 7 / 16)
26                blue = pixel[x + 1, y][2] + round(blue_error * 7 / 16)
27
28                pixel[x + 1, y] = (red, green, blue)
29
30            if x > 1 and y < y_lim - 1:
31                red = pixel[x - 1, y + 1][0] + round(red_error * 3 / 16)
32                green = pixel[x - 1, y + 1][1] + round(green_error * 3 / 16)
33                blue = pixel[x - 1, y + 1][2] + round(blue_error * 3 / 16)
34
35                pixel[x - 1, y + 1] = (red, green, blue)
36
37            if y < y_lim - 1:
38                red = pixel[x, y + 1][0] + round(red_error * 5 / 16)
39                green = pixel[x, y + 1][1] + round(green_error * 5 / 16)
40                blue = pixel[x, y + 1][2] + round(blue_error * 5 / 16)
41
42                pixel[x, y + 1] = (red, green, blue)
43
44            if x < x_lim - 1 and y < y_lim - 1:
45                red = pixel[x + 1, y + 1][0] + round(red_error * 1 / 16)
46                green = pixel[x + 1, y + 1][1] + round(green_error * 1 / 16)
47                blue = pixel[x + 1, y + 1][2] + round(blue_error * 1 / 16)
48
49                pixel[x + 1, y + 1] = (red, green, blue)
50
51    plt.imshow(new_img)
52    plt.show()
```



# Table of contents

- General introduction of computer vision
- Pixel Operations
- Quantizations
- Wrapping/ Combing / Scaling
- Image Convolution

- Moves a point to a new location by adding translation amounts to the coordinates of the point.



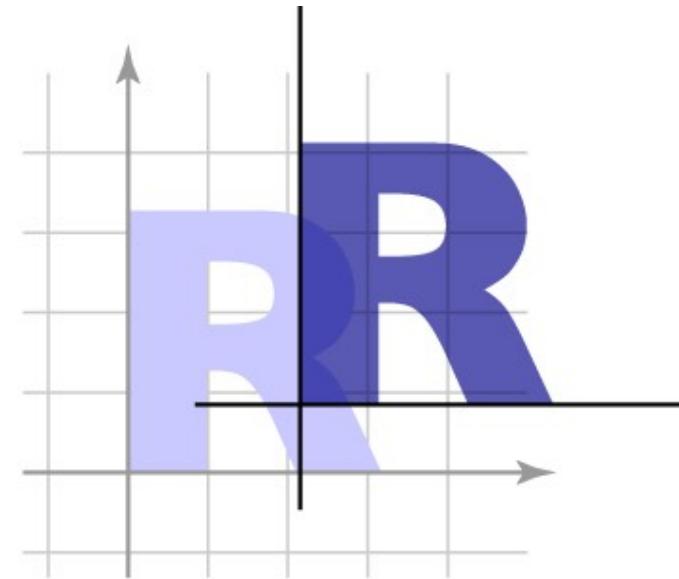
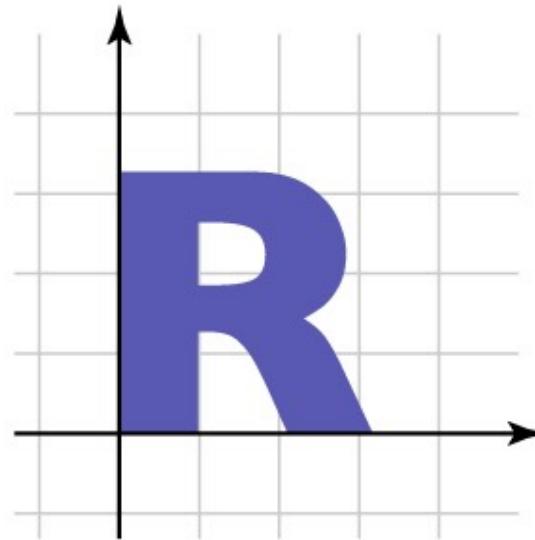
$$x' = x + dx, \quad y' = y + dy$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} dx \\ dy \end{bmatrix}$$

$$\underline{P' = P + T}$$

- Translation

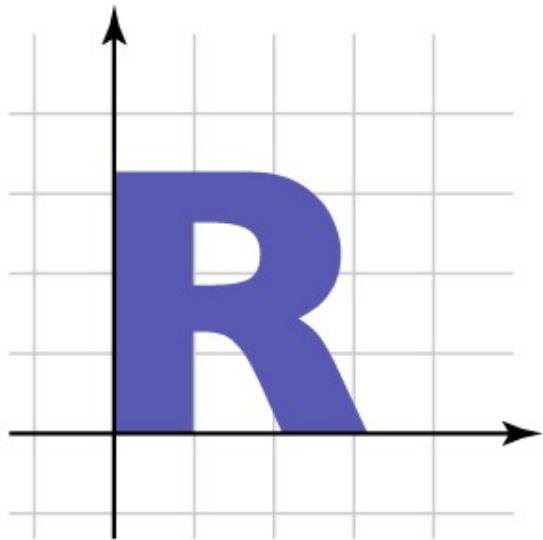
$$\begin{bmatrix} 1 & 0 & a_x \\ 0 & 1 & a_y \end{bmatrix} \quad \begin{bmatrix} 1 & 0 & 2.15 \\ 0 & 1 & 0.85 \end{bmatrix}$$



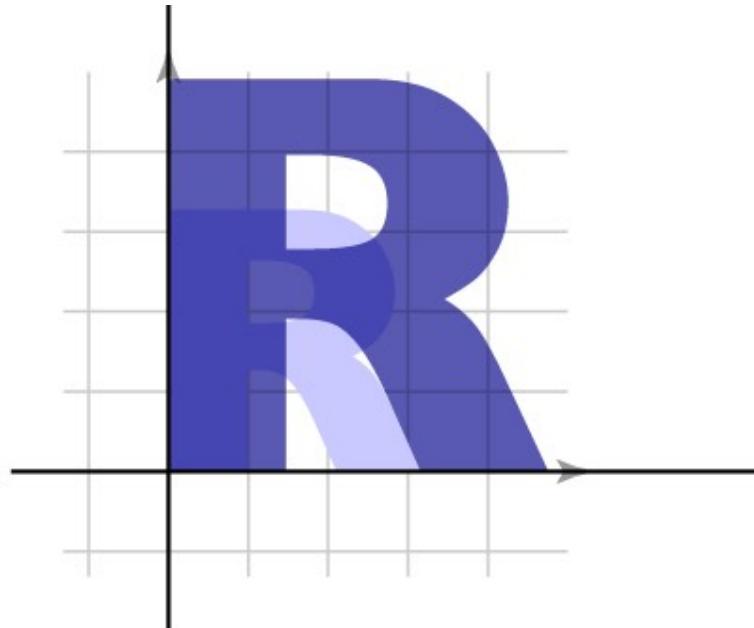
# Geometric operations - Uniform Scale



- Uniform scale  $\begin{bmatrix} s & 0 & a_x \\ 0 & s & a_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} sx + a_x \\ sy + a_y \\ 1 \end{bmatrix}$



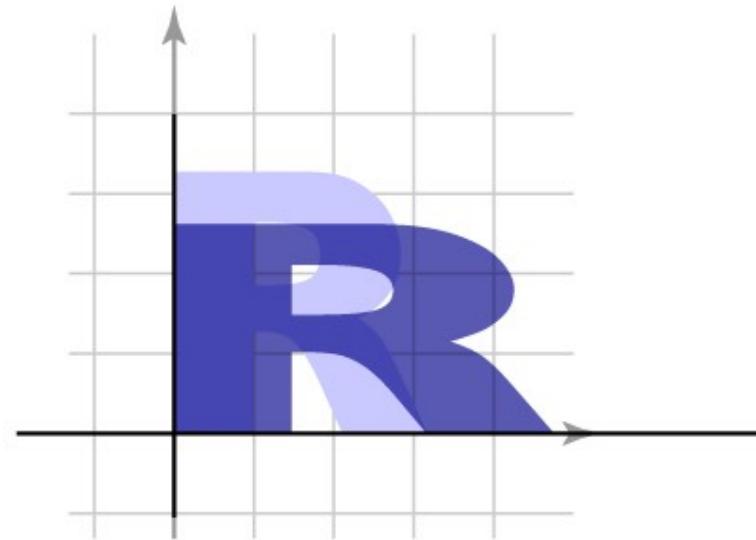
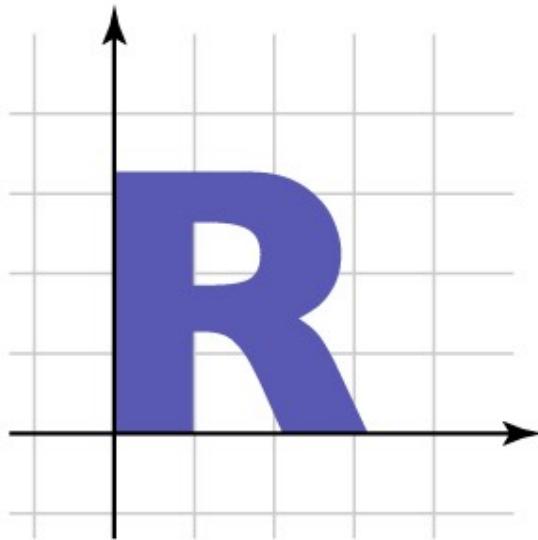
$$\begin{bmatrix} 1.5 & 0 & 0 \\ 0 & 1.5 & 0 \end{bmatrix}$$



- Nonuniform scale

$$\begin{bmatrix} s_x & 0 & a_x \\ 0 & s_y & a_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} s_x x + a_x \\ s_y y + a_y \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} 1.5 & 0 & 0 \\ 0 & 0.8 & 0 \end{bmatrix}$$

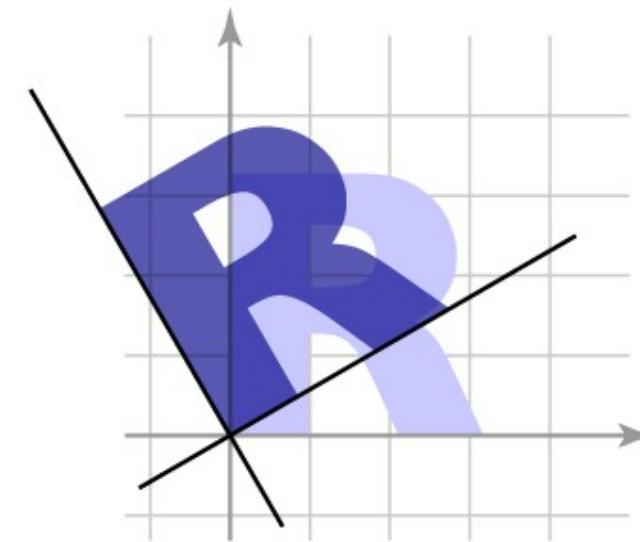
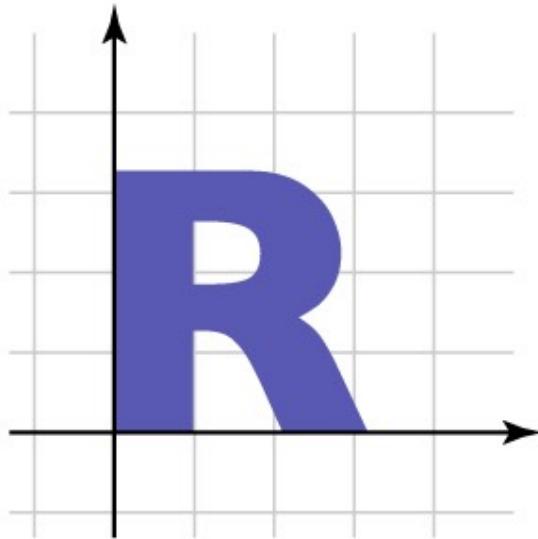


# Geometric operations - Rotation

- Rotation

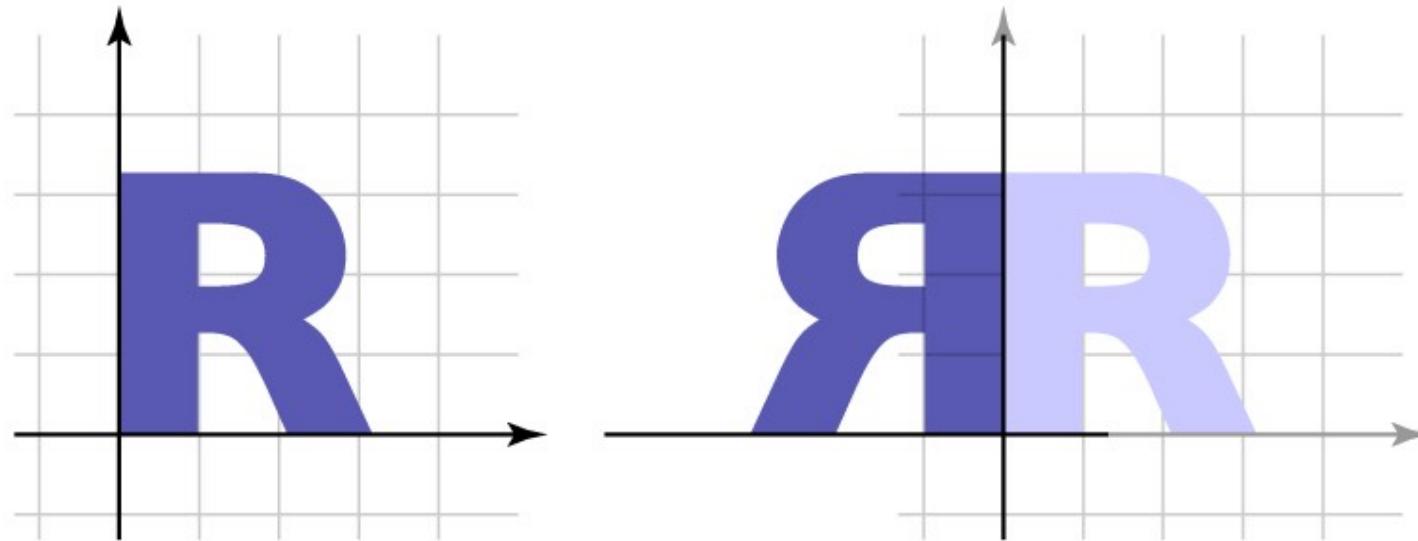
$$\begin{bmatrix} \cos \theta & -\sin \theta & a_x \\ \sin \theta & \cos \theta & a_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x \cos \theta - y \sin \theta + a_x \\ x \sin \theta + y \cos \theta + a_y \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} 0.86 & -0.5 & 0 \\ 0.5 & 0.866 & 0 \end{bmatrix}$$



- Reflection
  - can consider it a special case of nonuniform scale

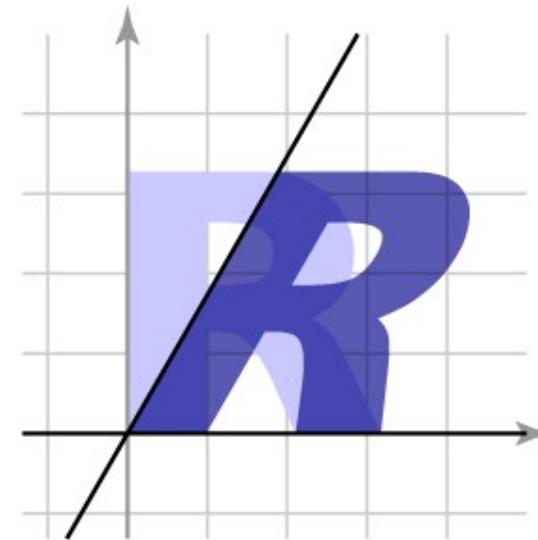
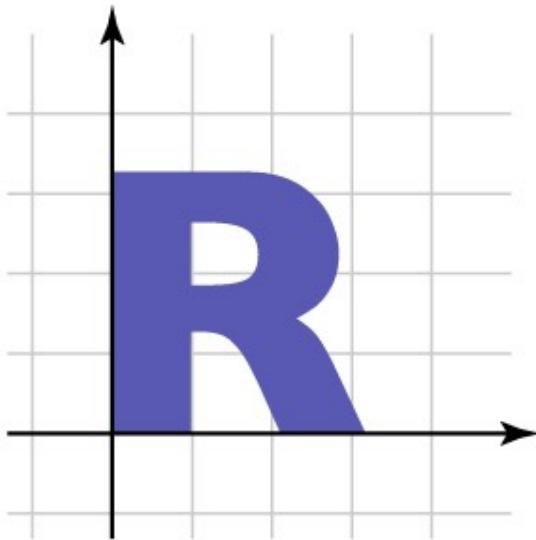
$$\begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$



# Geometric operations - Shear

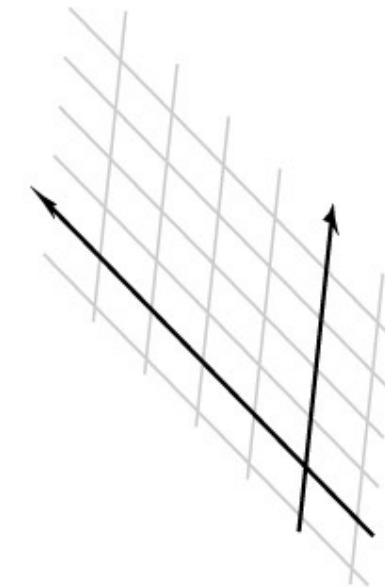
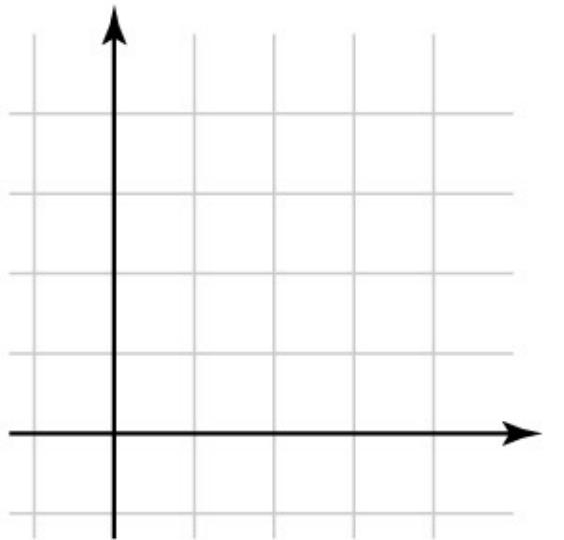
- Shear

$$\begin{bmatrix} 1 & a & a_x \\ 0 & 1 & a_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + ay + a_x \\ y + a_y \\ 1 \end{bmatrix}$$



# Affine transformations

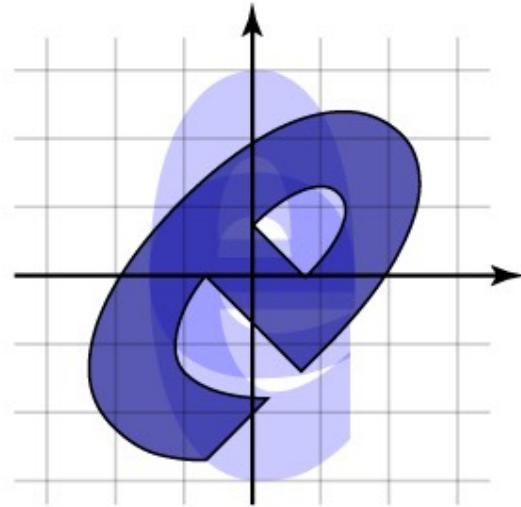
- The set of transformations we have been looking at is known as the “affine” transformations
  - straight lines preserved; parallel lines preserved
  - ratios of lengths along lines preserved (midpoints preserved)



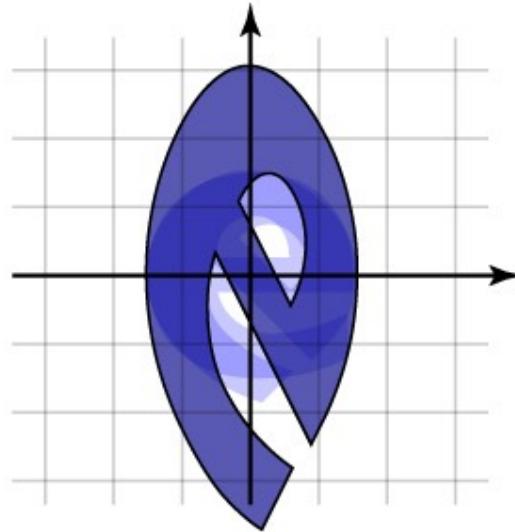
# Composite affine transformations

- Another example

$$\begin{bmatrix} s_{11} \cos \theta & -s_{12} \sin \theta & a_x \\ s_{21} \sin \theta & s_{22} \cos \theta & a_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} s_{11}x \cos \theta - s_{12}y \sin \theta + a_x \\ s_{21}x \sin \theta + s_{22}y \cos \theta + a_y \\ 1 \end{bmatrix}$$



scale, then rotate

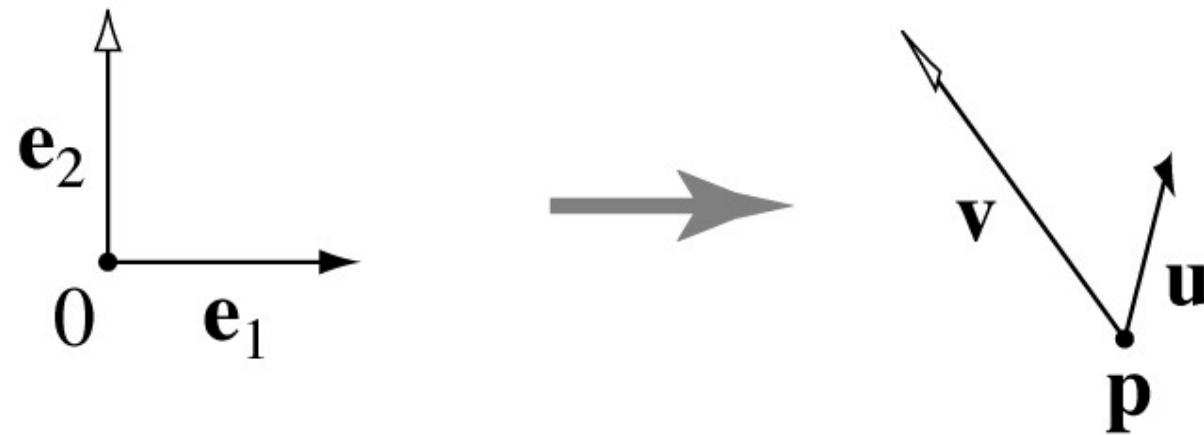


rotate, then scale

# Affine change of coordinates

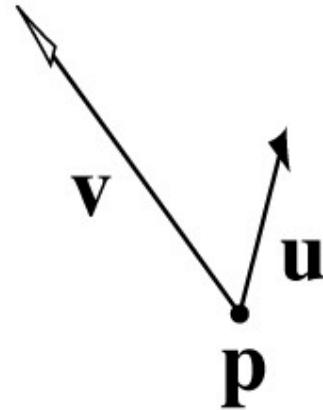
- Six degrees of freedom

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \end{bmatrix} \quad \text{or} \quad [u \quad v \quad p]$$



# Affine change of coordinates

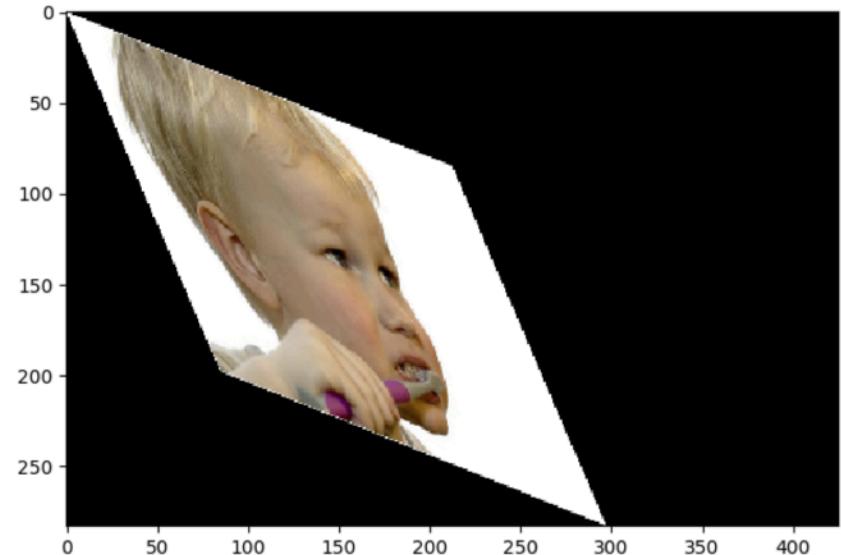
- Coordinate frame: point plus basis
- Interpretation: transformation changes representation of point from one basis to another
- “Frame to canonical” matrix has frame in columns
  - takes points represented in frame
  - represents them in canonical basis
  - e.g.  $[0\ 0]$ ,  $[1\ 0]$ ,  $[0\ 1]$
- Seems backward but bears thinking about



$$\begin{bmatrix} u & v & p \end{bmatrix}$$

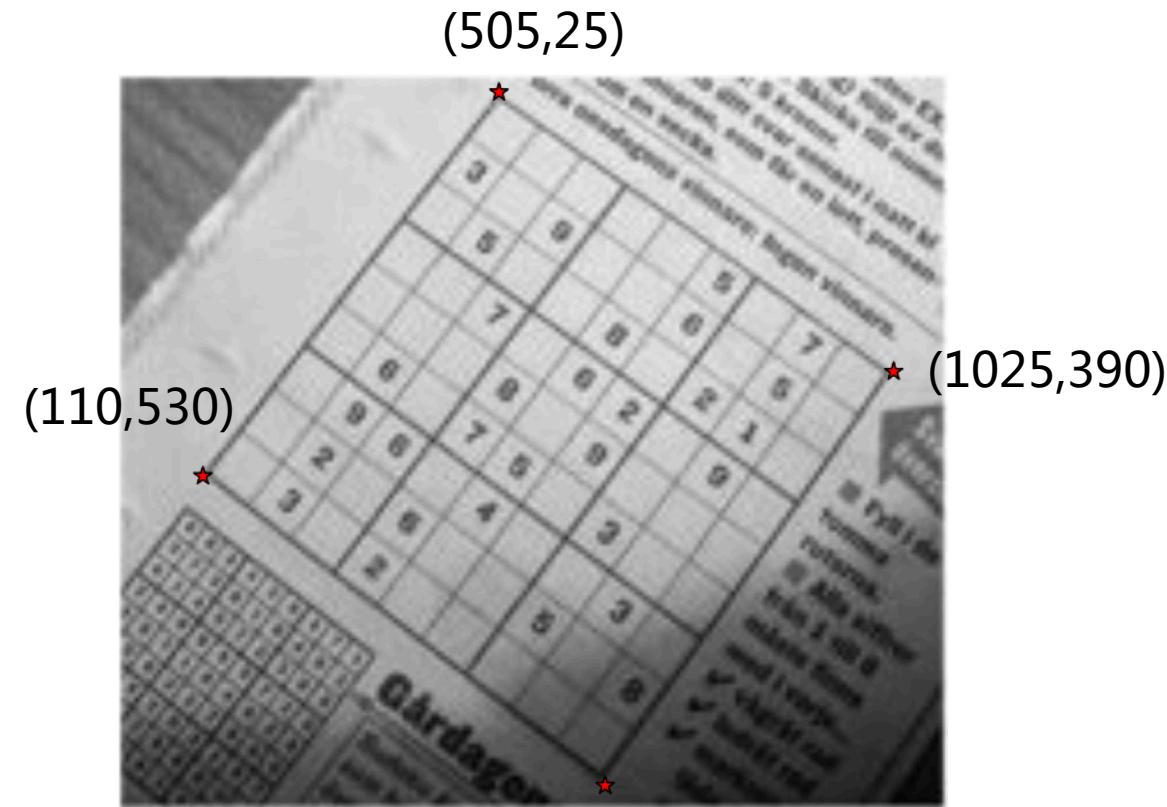
# Affine change of coordinates

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4 img = cv2.imread('ChildBrushingTeeth.jpg',cv2.IMREAD_COLOR)
5 img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
6 (height,width,deep) = img.shape
7 matRotation = np.array(
8     [[0.5,0.3,0],
9      [0.2, 0.7, 0]])
10 dest_img= cv2.warpAffine(img,matRotation,(0,0))
11 plt.imshow(dest_img)
12 plt.figure(1)
13 plt.show()
14
```

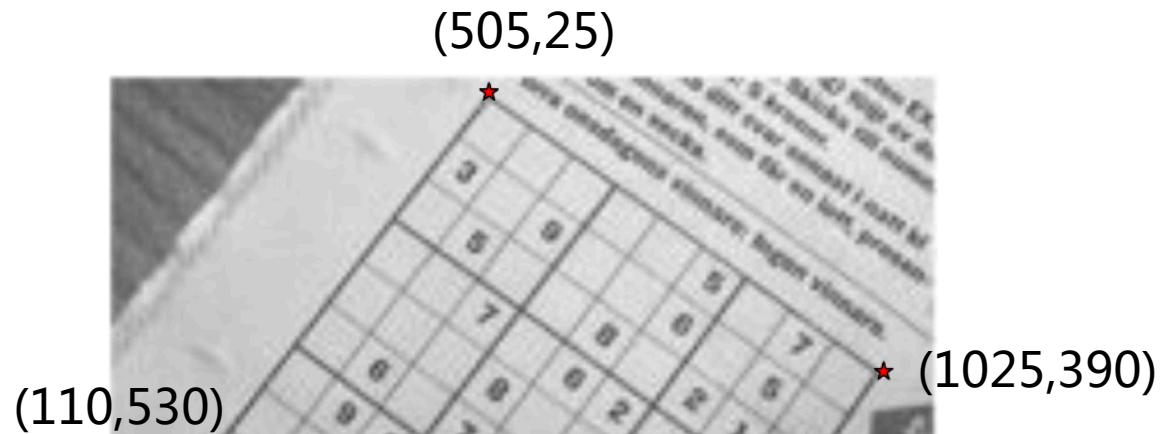


# Affine change of coordinates

```
1 import cv2
2 import numpy as np
3 img = cv2.imread('sudoku.png',cv2.IMREAD_COLOR)
4 (height,width,deep) = img.shape
5 # lefttop, letbottom,rightbutton
6 matSrc = np.float32([[505,25],[110,530],[1025,390]])
7 matDest = np.float32([[0,0],[0, 970],[1100,0]])
8 matAffine = cv2.getAffineTransform(matSrc,matDest)
9 dest_img = cv2.warpAffine(img,matAffine,(0,0))
10 cv2.imshow('img',dest_img)
11 cv2.waitKey(0)
```



# Affine change of coordinates



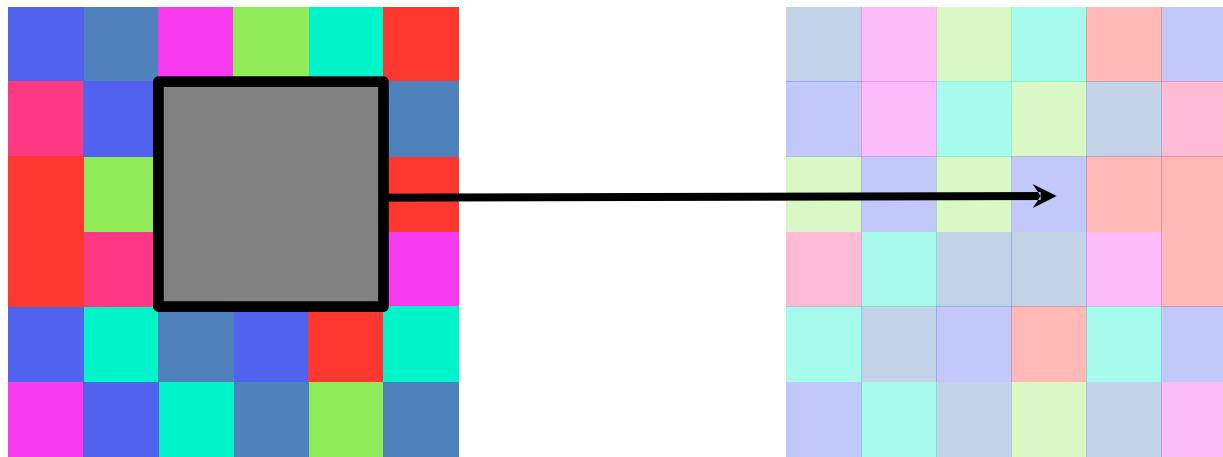
(0,0) (1100,0)



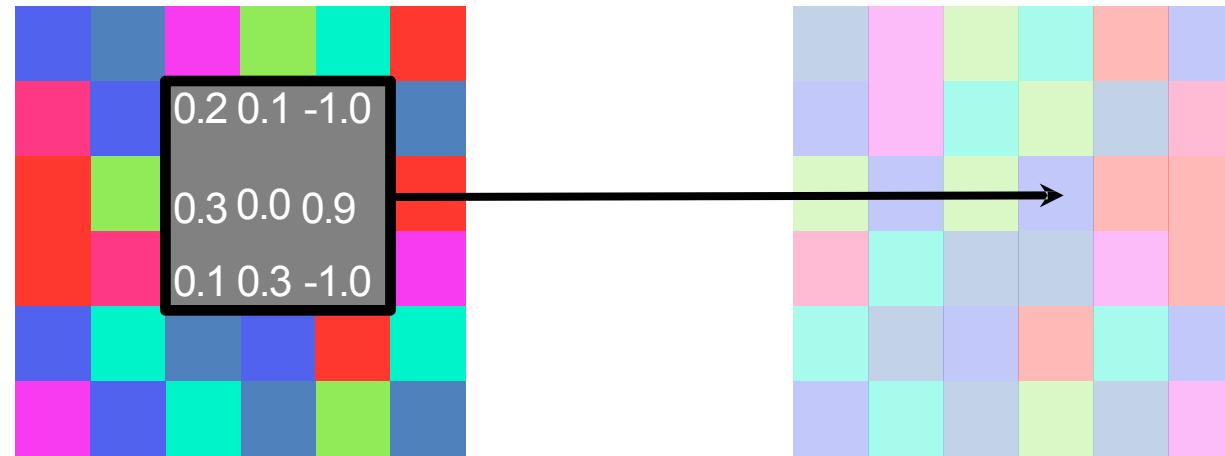
# Table of contents

- General introduction of computer vision
- Pixel Operations
- Quantizations
- Wrapping/ Combing / Scaling
- Image Convolution
  - Image convolution overview
  - Image derivate and Sobel Kernels
  - Laplacian Kernels
  - Image smoothing and mean kernels
  - Gaussian Kernels
  - Edge detection (canny algorithm)

# Convolution: Neighborhood Operations

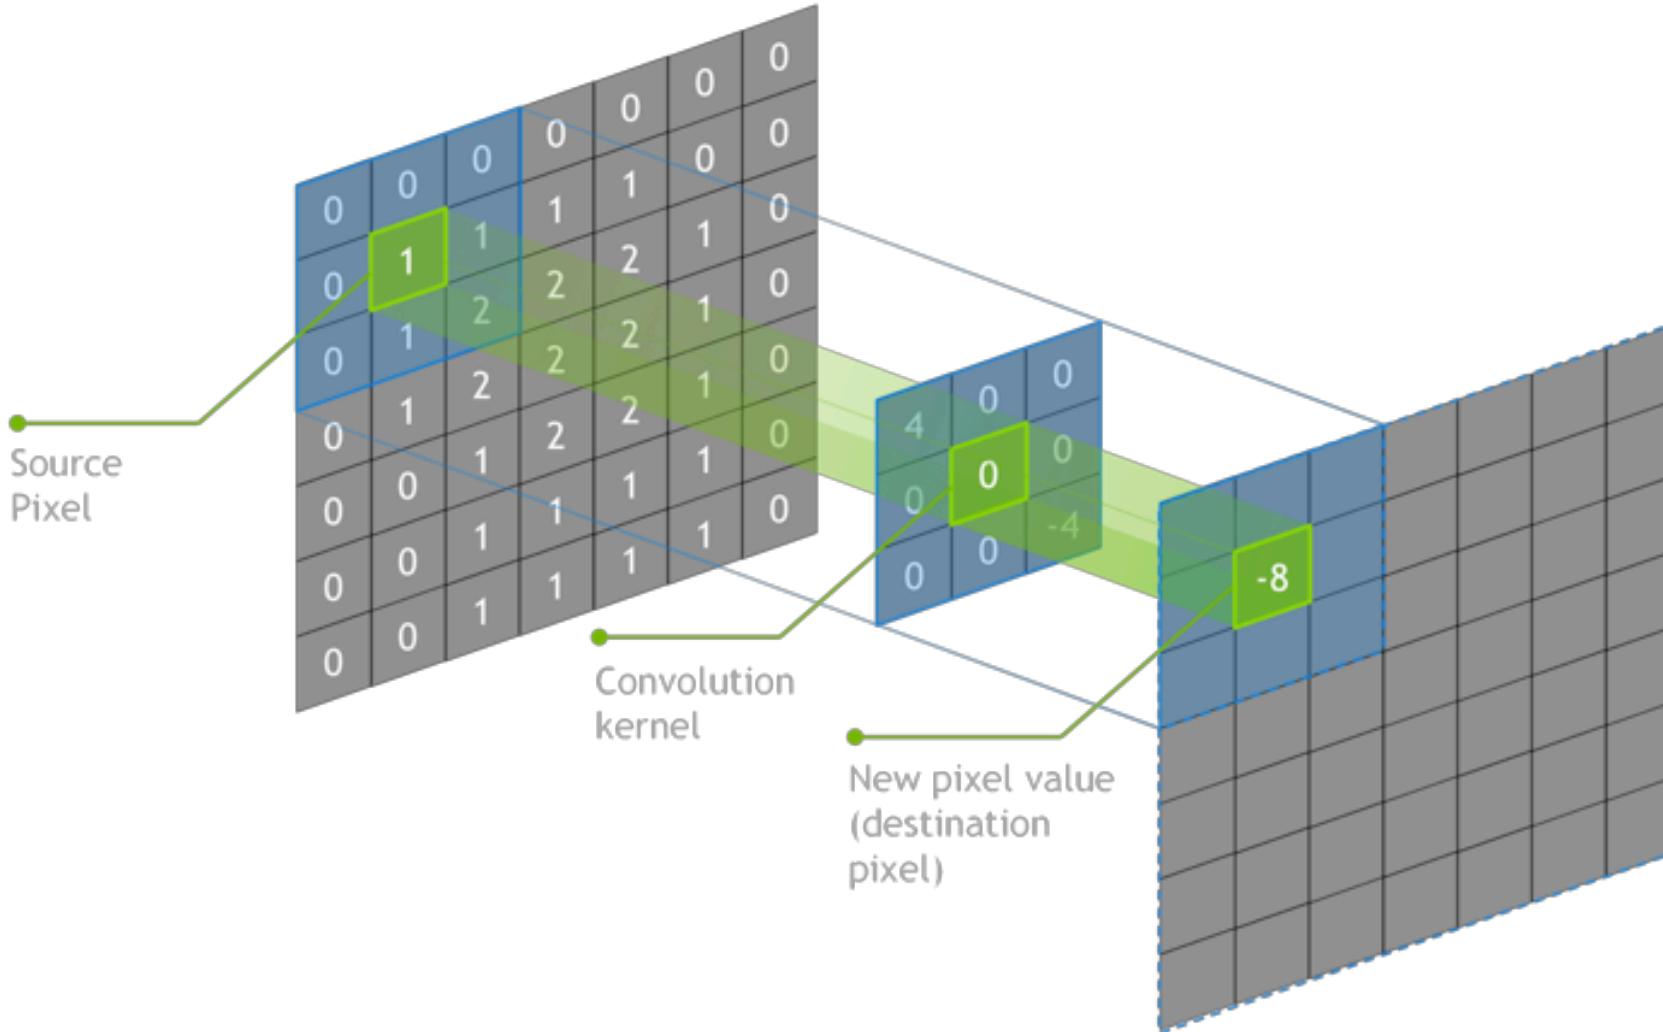


# Convolution



$$F = \begin{bmatrix} 0.2 & 0.1 & -1.0 \\ 0.3 & 0.0 & 0.9 \\ 0.1 & 0.3 & -1.0 \end{bmatrix} \quad I' = F * I$$

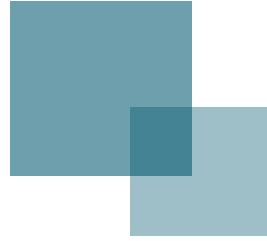
# Convolution images



# Convolution Kernels

Operation	Kernel $\omega$	Image result $g(x,y)$
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	

Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur $3 \times 3$ (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	
Gaussian blur $5 \times 5$ (approximation)	$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	
Unsharp masking $5 \times 5$ Based on Gaussian blur with amount as 1 and threshold as 0 (with no image mask)	$\frac{-1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & -476 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	



# Thanks

FAQ时间